

Vještačka inteligencija

Predavanje 3: Klasifikacija

*Jedinu pravu sigurnost u današnjem svijetu
čovjeku mogu pružiti znanje, iskustvo i
sposobnost.
~Henry Ford*

Odgovorna nastavnica: Vanr. prof. dr Amila Akagić

Univerzitet u Sarajevu



Uvodne informacije

- This work is licensed under a Creative Commons 'Attribution-NonCommercial-ShareAlike 4.0 International' license. EN: <https://creativecommons.org/licenses/by-nc-sa/4.0/>



- Ovaj rad je licenciran pod medunarodnom licencom 'Imenovanje-Nekomercijalno-Dijeli pod istim uvjetima 4.0' od strane Creative Commons. HR:

<https://creativecommons.org/licenses/by-nc-sa/4.0/deed.hr>

Najave

- ❑ Danas u 14:00 je prvi kviz u okviru realizacije vježbi.
 - ❑ 5 pitanja, 5 minuta vremena.
 - ❑ Automatsko ocjenjivanje nakon isteka vremena.

Aktuelnosti / Vještačka inteligencija

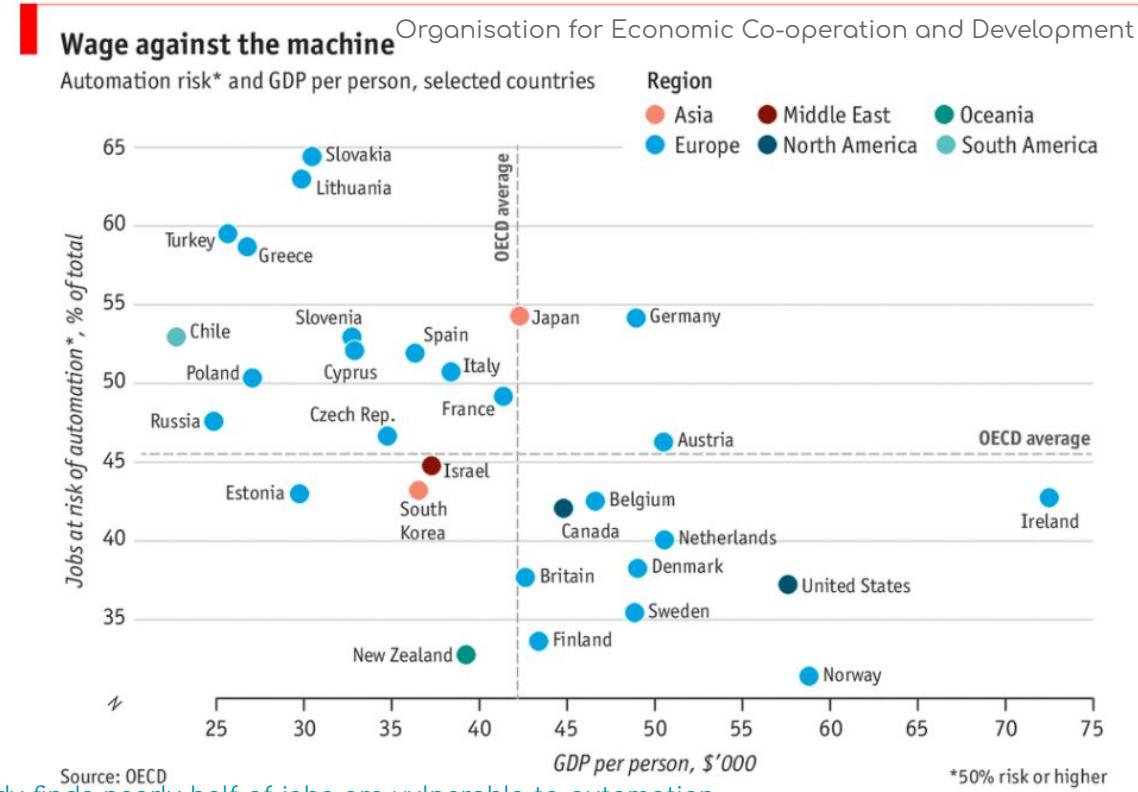
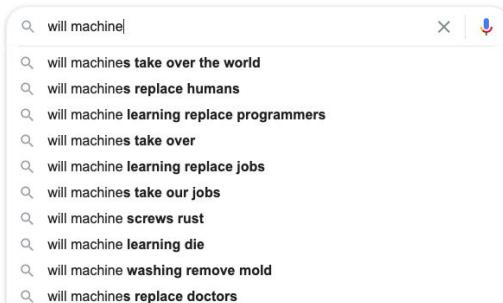
- Da li će VI (AI) ugrožavati / ugrožava radna mjesta? Oblast pružanja usluga?

Daily chart

A study finds nearly half of jobs are vulnerable to automation

That could free people to pursue more interesting careers

Google



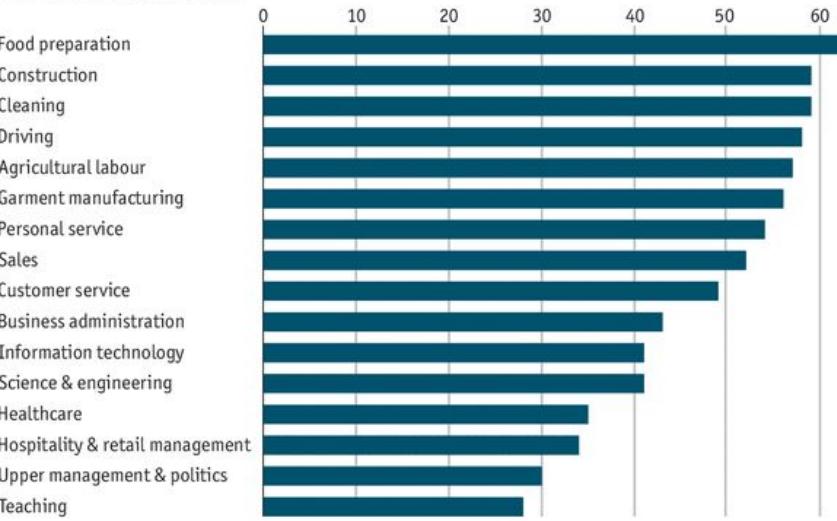
<https://www.economist.com/graphic-detail/2018/04/24/a-study-finds-nearly-half-of-jobs-are-vulnerable-to-automation>

Aktuelnosti / Vještačka inteligencija

- ❑ U 2013. godini, Carl Benedikt Frey i Michael Osborne sa Univerziteta Oxford koristili su algoritam mašinskog učenja kako bi procijenili koliko lagano se može automatizirati ~700 različitih vrsta poslova u Americi.
 - ❑ Zaključili su da bi računari (mašine) u potpunosti mogle izvršavati 47% poslova u narednih deset ili dvadeset godina.
 - ❑ Koji su to poslovi?
 - ❑ Poslovi u Slovačkoj su dva puta više ranjiviji nego poslovi u Norveškoj.
 - ❑ Generalno, poslovi u bogatijim državama su znatno manje ranjiviji.
 - ❑ Na primjer, u Južnoj Koreji 30% radnih mesta je u prerađivačkoj industriji, u usporedbi s 22% u Kanadi. Bez obzira na to, u prosjeku je teže automatizirati korejske poslove nego kanadske (zbog vrste posla).

Automated for the people

Automation risk by job type, %



In April, the OECD criticised an influential 2013 forecast by Oxford University that found about 47% of jobs in the US in 2010 and 35% in the UK were at "high risk" of being automated over the following 20 years.

The OECD instead put the US figure at about 10% and the UK's at 12% - although it did suggest many more workers would see their tasks changing significantly.

Aktuelnosti / Vještačka inteligencija

Quora

Q Search for questions, people, and topics

Software Development Software Developers Computer Programmers Computer Science

Software Engineering Software and Applications Computer Programming

Will GPT-3 replace software developers?

3 Answers



Alastair Stell, Javascript, Python, C++, C#, Smalltalk, C, Web Programming

Answered July 25, 2020 · Author has 440 answers and 649.1K answer views

Coders? Yes. Software engineers? Not so much. GPT-3 rather like Watson, is essentially focused on machine learning. It doesn't have creative thought and problem solving is restricted to "what is known" and what can be projected from data. That may sound like AI but it isn't because GPT-3 doesn't understand what it is doing - there is no conscious process involved - which severely limits creativity.

People also ask

Will GPT 3 kill coding?

GPT-3 would never kill jobs skilled developers. Instead its a wake-up call for cargo coders and developers. It'll urge them to buckle up and upskill to ensure they're up for solving complex computer **programming** problems.

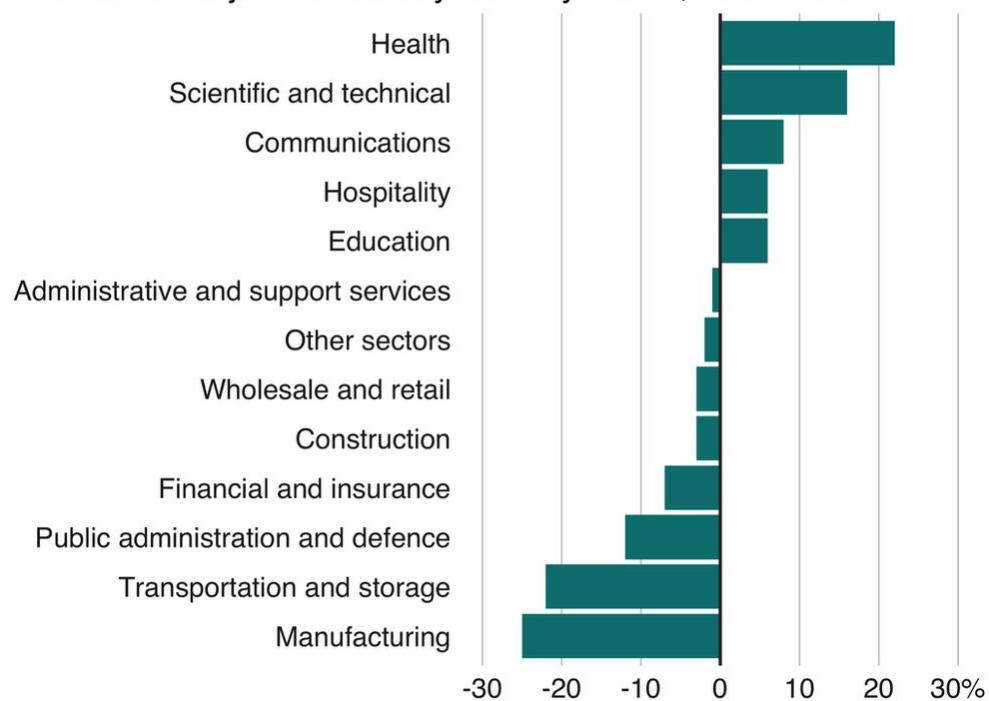
Aktuelnosti / Vještačka inteligencija

- ❑ AI će kreirati onoliko poslova koliko i uništi?

"Vjerovatno će četvrta industrijska revolucija favorizirati one s jakim digitalnim vještinama, kao i sposobnostima poput kreativnosti i timskog rada koje je mašinama teže replicirati."

How AI could change the job market

Estimated net job creation by industry sector, 2017-2037



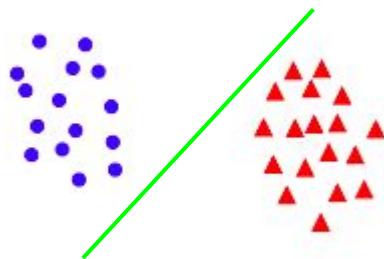
Source: PwC

BBC

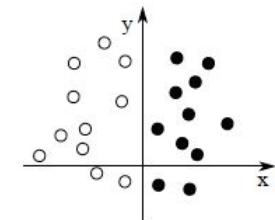
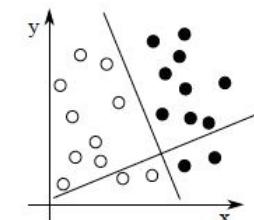
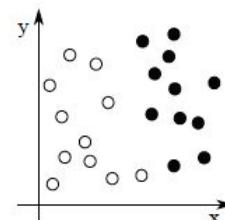
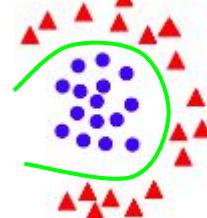
Klasifikator

- ❑ **Klasifikator** je bilo koji algoritam koji sortira podatke u označene klase ili kategorije informacija.
 - ❑ Primjeri:
 - ❑ Filteri za neželjenu poštu koji skeniraju dolazne email-ove i klasificiraju ih kao "neželjenu poštu" (*spam*) ili "poželjenu poštu".
 - ❑ Predikcija da li osoba ima neko oboljenje na osnovu predočenih nalaza.
 - ❑ Predikcija da li je neka finansijska transakcija lažna (*fraud detection*).
- ❑ Klasifikatori su konkretna implementacija prepoznavanja obrazaca (*pattern recognition*) u mnogim oblicima mašinskog učenja.

Prvi skup podataka

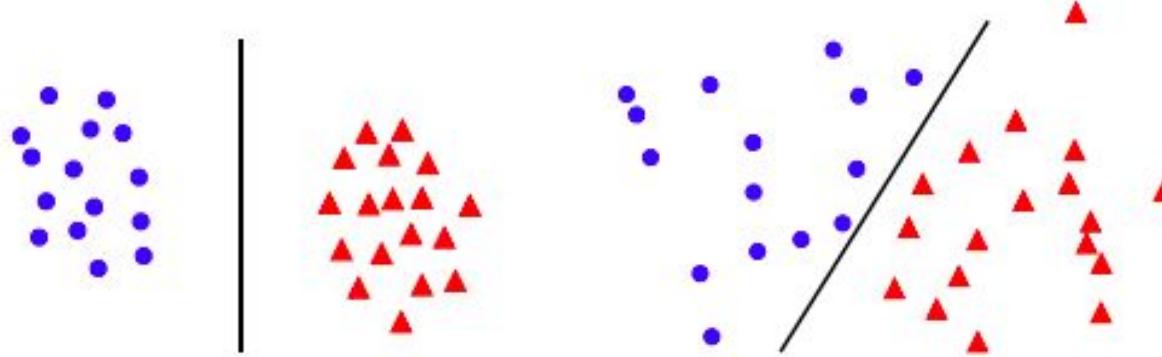


Drugi skup podataka

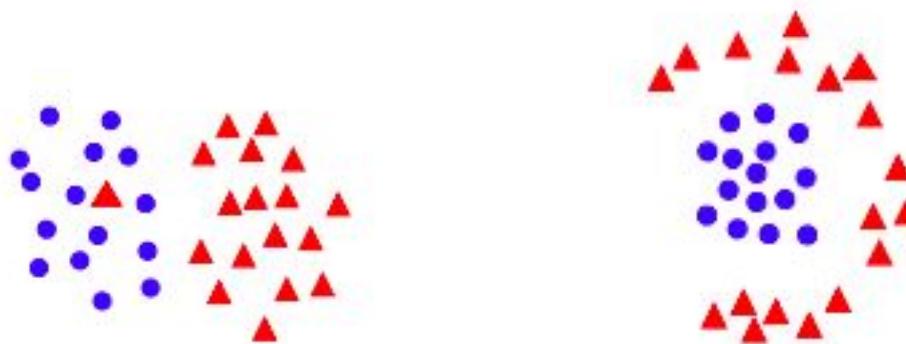


Primjeri skupova

Mogu se klasificirati linearnim klasifikatorom



Ne mogu se klasificirati linearnim klasifikatorom



Klasifikator

- ❑ Primjeri klasifikatora iz prakse:
 - ❑ Perceptron
 - ❑ Naivni Bayes
 - ❑ Stablo odlučivanja
 - ❑ Logistička regresija
 - ❑ K-Nearest Neighbor
 - ❑ Vještačke neuronske mreže / duboko učenje
 - ❑ Support Vector Machine (SVM)
 - ❑

1958: Perceptron

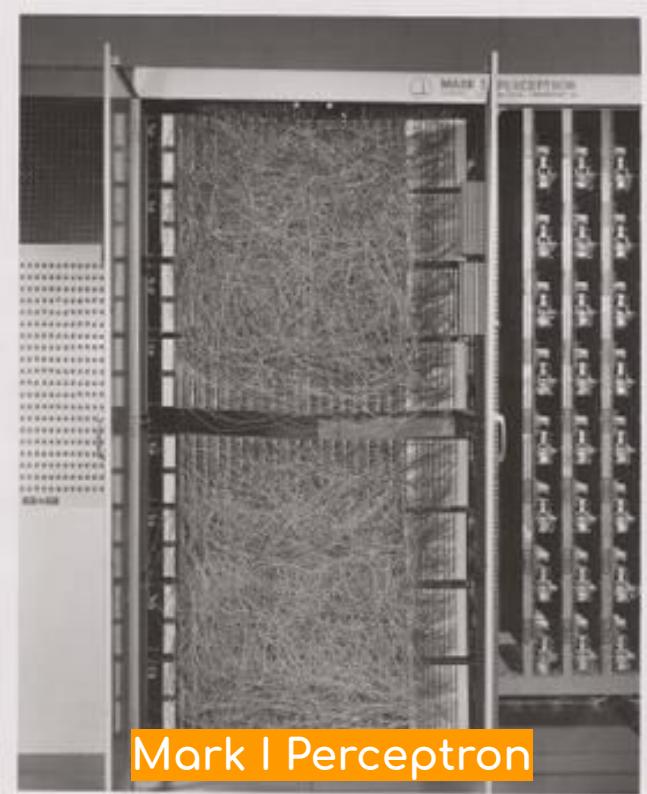
- ❑ Jedan od prvih algoritama koji je imao mogućnost učenja iz podataka.
- ❑ Bio je implementiran u hardveru!
 - ❑ Vrijednosti koeficijenta su čuvani u potenciometarima, a električni motori su korišteni za njihovo ažuriranje!!!
- ❑ Rad perceptron-a je demonstriran na prepoznavanju alfabeta sa 20x20 slike.
Kamera je koristila fotoćelije kadmijum sulfida.
- ❑ Po današnjoj klasifikaciji metoda, ova metoda bi spadala u metode linearnih klasifikatora.



1958
Perceptron



Frank Rosenblatt



Mark I Perceptron

1958: Perceptron

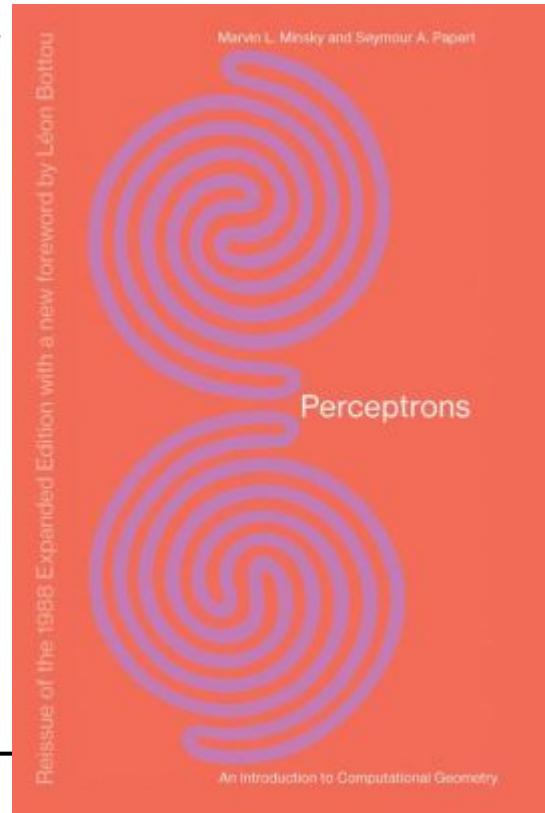
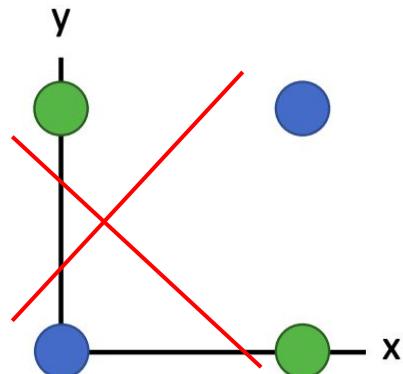


"The Navy revealed the embryo of an electronic computer today that it expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence." New York Times Archives, July 8, 1958, Page 25

1969: Knjiga “Perceptrons...” (Minsky & Papert)

- 1969: “*Perceptrons: an introduction to computational geometry*”, knjiga autora Marvin Minsky i Seymour Papert.
- Matematički dokaz o ograničenjima dvoslojnih *feed forward* perceptrona.

X	Y	F(x,y)
0	0	0
0	1	1
1	0	1
1	1	0

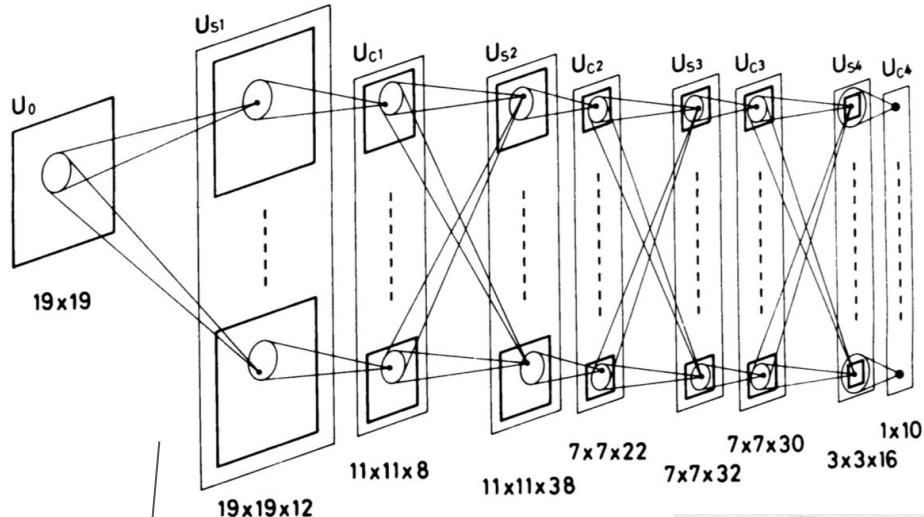


1958
Perceptron

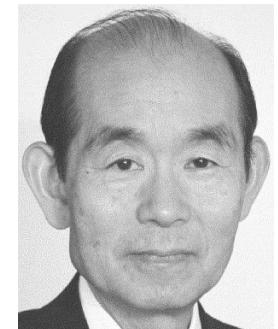
1969
Minsky & Papert

1980: Neocognitron

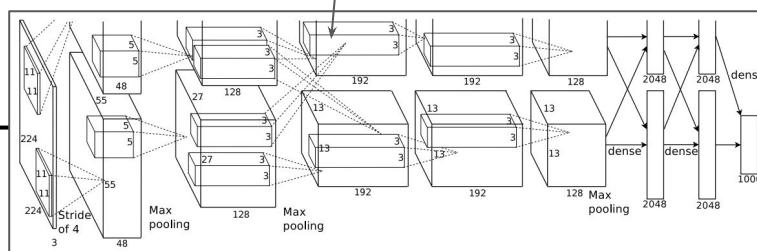
- 1980: Novi algoritam za prepoznavanje šablonu (*pattern recognition*) baziran na modelu vizuelnog sistema.
- Višeslojna neuronska mreža.
- Neocognitron je samoorganizirajuća mreža nastala učenjem bez nadzora.
- Osnovne karakteristike: hijerarhija zasnovana na jednostavnim (**konvolucija**) i složenim (**pooling**) ćelijama.



32 godine!



Kunihiko Fukushima



1958
Perceptron

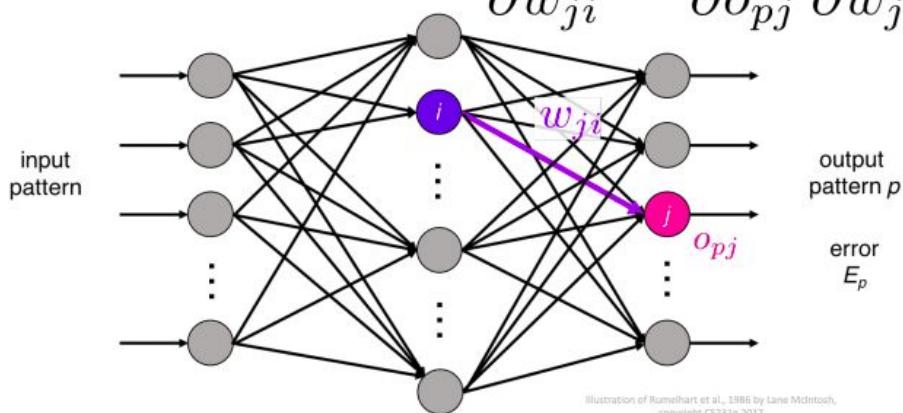
1969
Minsky & Papert

1980
Fukushima

1986: Backprop / Rumelhart, Hinton i Williams

- Uvođenjem backpropagacije omogućeno je uspješno treniranje višeslojnih perceptrona.
- Koristi se kod izračunavanja gradijenta u neuronskim mrežama.

$$\frac{\partial E_p}{\partial w_{ji}} = \frac{\partial E_p}{\partial o_{pj}} \frac{\partial o_{pj}}{\partial w_{ji}}$$



Learning representations by back-propagating errors

David E. Rumelhart*, Geoffrey E. Hinton†
& Ronald J. Williams*

* Institute for Cognitive Science, C-015, University of California,
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,
Pittsburgh, Philadelphia 15213, USA



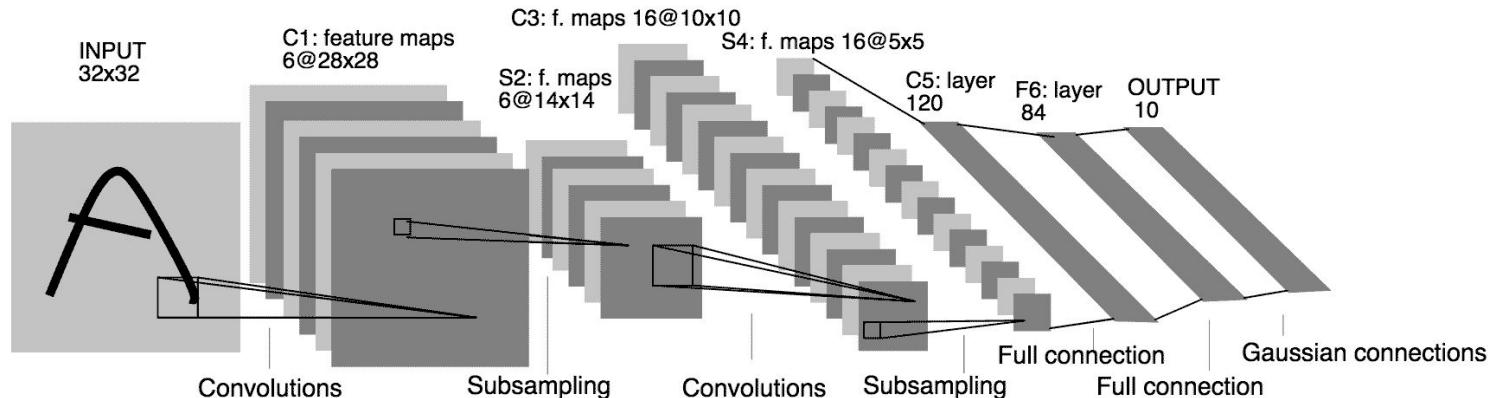
1958
Perceptron

1969
Minsky & Papert

1980
Fukushima

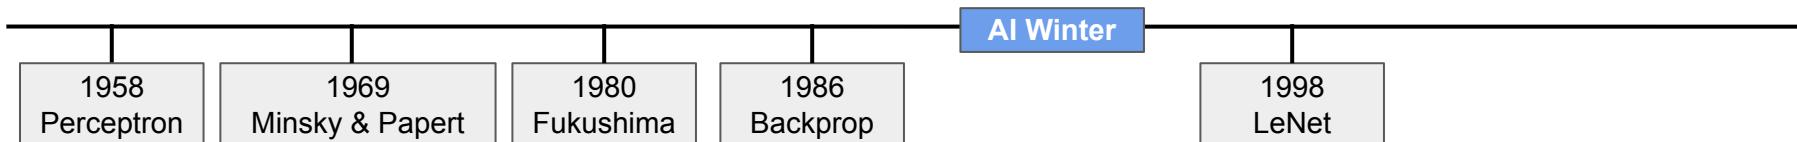
1986
Backprop

1998: Konvolucijske mreže (LeCun)



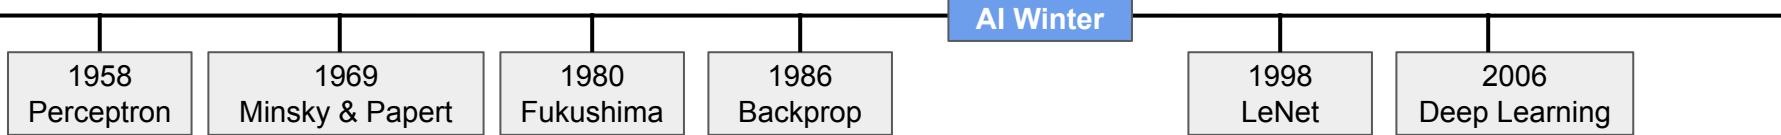
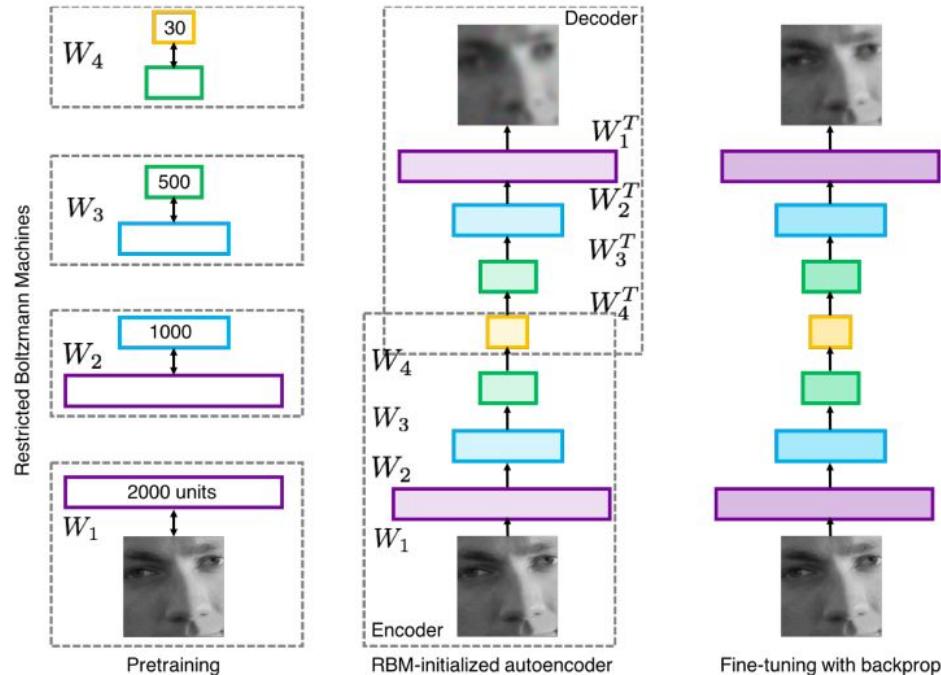
- Prvi put primijenjena backpropagacija na mrežu koja podsjeća na Neocognitron.
- Primjena: prepoznavanje rukom pisanih slova alfabeta.
- Sistem je korišten od strane kompanije NEC za prepoznavanje rukom pisanih tekstova.

Source: LeCun, Yann, et al. "Gradient-based learning applied to document recognition." Proceedings of the IEEE 86.11 (1998) 2278-2324.



2000-te: Duboko učenje ili Deep Learning (DL)

- Osnovna ideja: nakon uspjeha LeNet-a istraživači su pokušali trenirati sve dublje i dublje mreže (sa više slojeva).
- Trebalo je značajno vremena da ova ideja postane “mainstream”.
- Značajni radovi iz oblasti:
 - Hinton and Salakhutdinov, 2006
 - Bengio et al, 2007
 - Lee et al, 2009
 - Glorot and Bengio, 2010



2009: ImageNet takmičenje



1,000 object classes
1,431,167 images



Primjer izlaza:
Cat
Dog
Car
Mouse
Hat

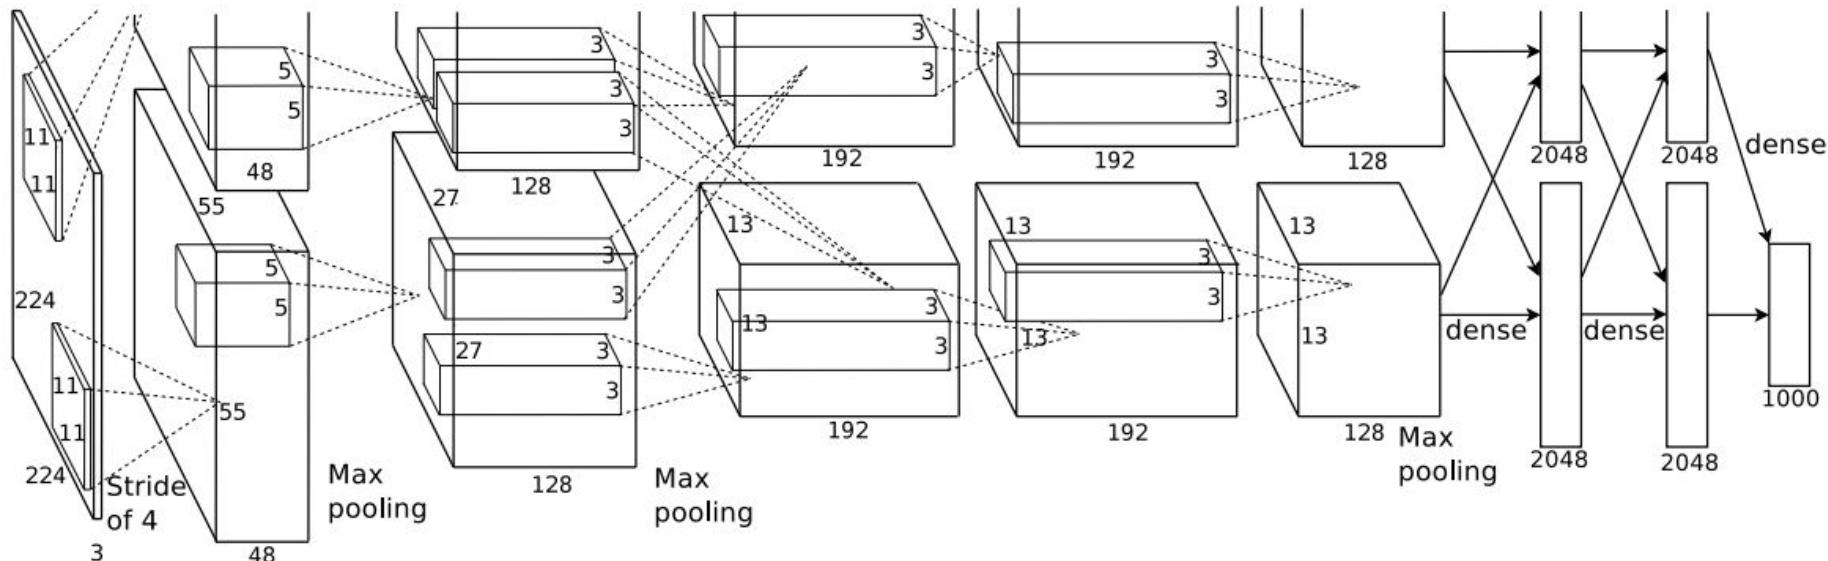
2012: AlexNet

ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky
University of Toronto

Ilya Sutskever
University of Toronto

Geoffrey E. Hinton
University of Toronto



ImageNet Classification with Deep Convolutional Neural Networks, Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton
<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>

2012: Primjeri korištenja ConvNets

Image Classification

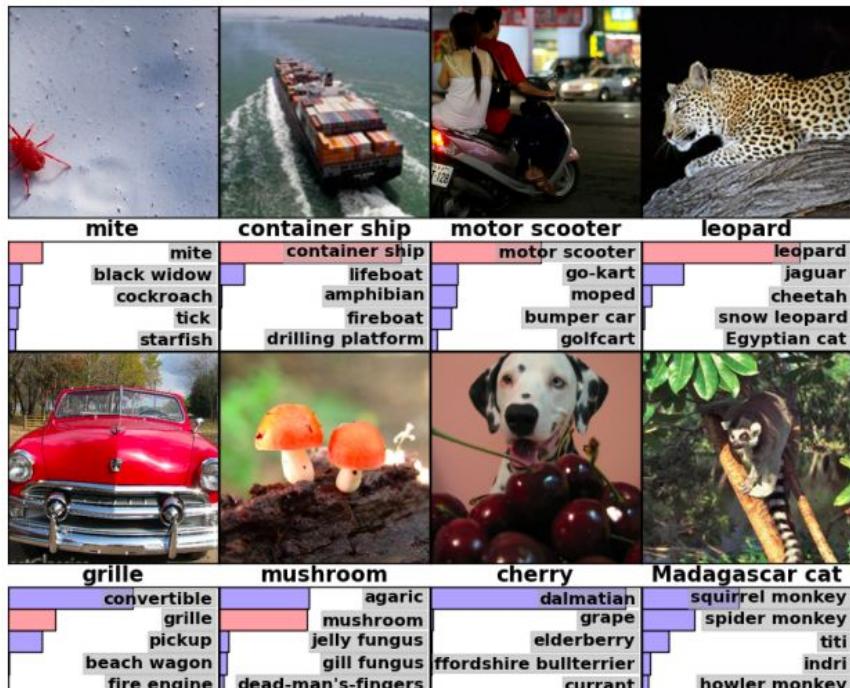
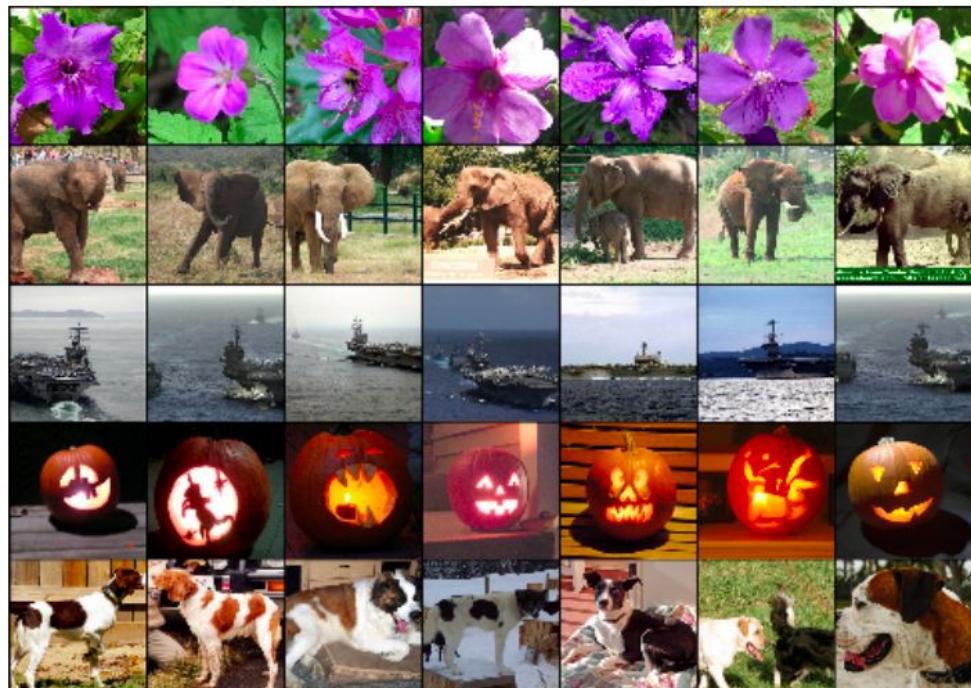


Image Retrieval



ImageNet Classification with Deep Convolutional Neural Networks, Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton

<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>

Primjeri korištenja ConvNets

Object Detection

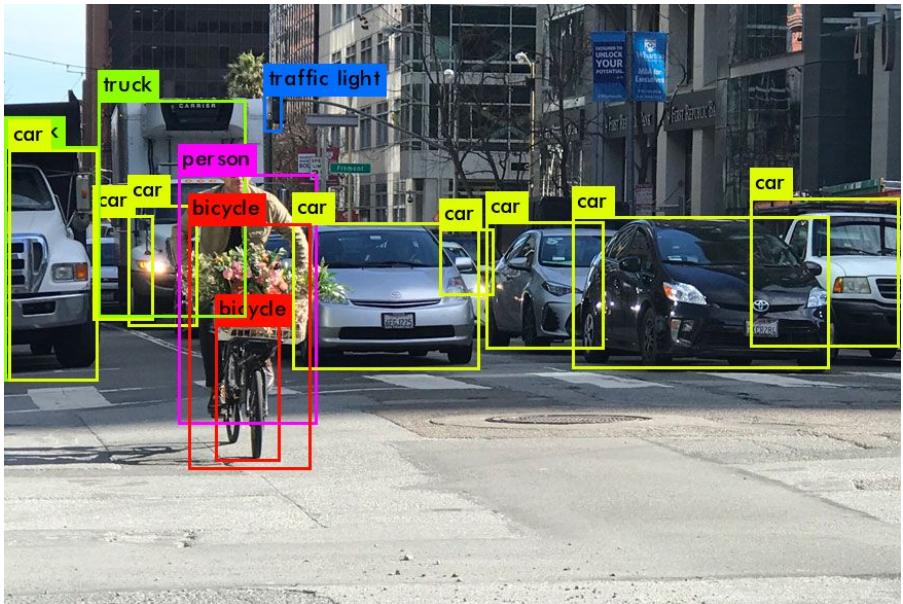


Image Segmentation



Road	Sidewalk	Building	Fence
Pole	Vegetation	Vehicle	Unlabel

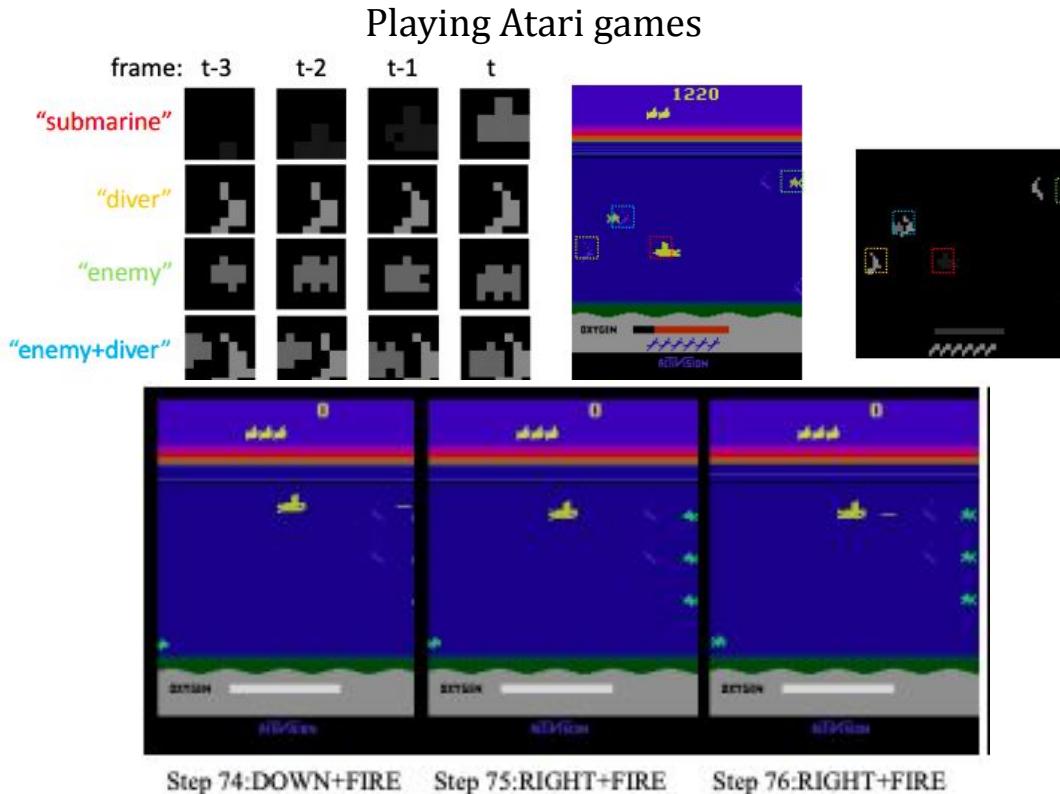
Primjeri korištenja ConvNets

Pose Recognition



DeepPose: Human Pose Estimation via Deep Neural Networks, Alexander Toshev, Christian Szegedy (Google, 2014)
<https://arxiv.org/pdf/1312.4659.pdf>

Sistem za igranje igrica?



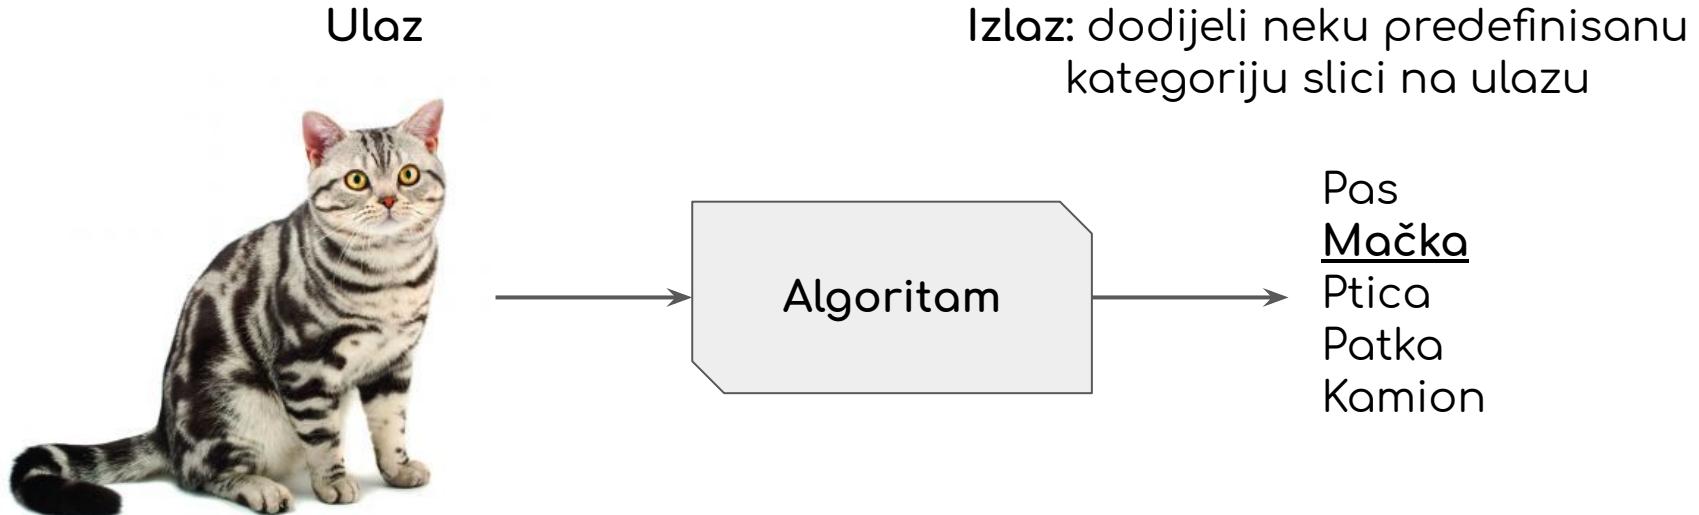
Deep Learning for Real-Time Atari Game Play Using Offline Monte-Carlo Tree Search Planning (2014)
<https://papers.nips.cc/paper/5421-deep-learning-for-real-time-atari-game-play-using-offline-monte-carlo-tree-search-planning.pdf>

Kako koristiti podatke za proces učenja?

- ❑ Prošli put smo vidjeli različite vrste podataka koje koristimo za treniranje različitih modela za učenje.
- ❑ Koji je najbolji način da ih koristimo? Koje su to ustaljene prakse podjele jednog skupa podataka?
- ❑ U slučaju KNN, kako izabrati K i $distancu$?
 - ❑ K i $distanca$ su primjeri hiperparametara koje biramo prije treniranja!
 - ❑ Ove parametre ne učimo iz podataka, nego ih *a priori* određujemo.
 - ❑ Izbor ovih parametara će uveliko zavisiti od problema koji pokušavamo riješiti.
 - ❑ Najbolje rješenje: Isprobati sve varijacije i odlučiti se za najbolju.

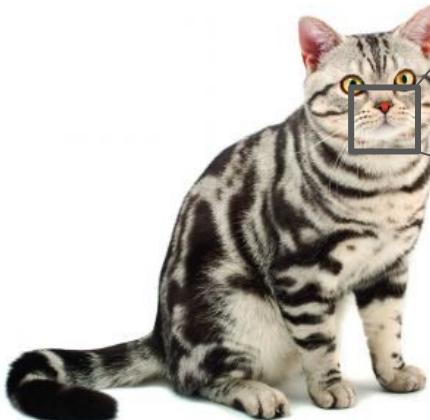
Klasifikacija slika

Klasifikacija slika: osnovni koncept



Za čovjeka je klasifikacija slika trivijalan zadatak, međutim za računar je to složen zadatak prvenstveno zbog toga što računar slike vidi u drugačijem obliku (nema jedinstveni vizuelni sistem).

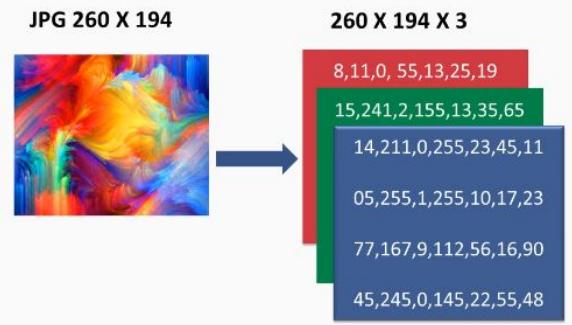
Predstavljanje slika u računaru



```
[ (72, 89, 83), (66, 81, 76), (58, 73, 70), (53, 65, 63), (42, 54, 52), (45, 55, 54), (44, 55, 51), (40, 49, 46), (36, 45, 40), (40, 47, 40), (48, 55, 48), (55, 62, 54), (50, 53, 44), (66, 69, 60), (109, 112, 103), (122, 125, 116), (124, 127, 116), (119, 122, 111), (99, 101, 88), (109, 111, 98), (124, 126, 112), (137, 139, 125), (142, 145, 128), (124, 125, 109), (120, 121, 103), (133, 134, 116), (141, 142, 124), (159, 160, 142), (165, 159, 143), (154, 148, 132), (168, 162, 146), (160, 154, 138), (141, 136, 117), (151, 144, 126), (171, 164, 146), (190, 183, 165), (187, 180, 161), (193, 185, 166), (195, 187, 168), (195, 187, 168), (199, 191, 172), (206, 198, 179), (210, 202, 181), (210, 202, 181), (208, 200, 179), (197, 189, 168), (181, 170, 150), (158, 147, 127)]
```

mačka?

Šta računar vidi?
Šta ako se promijeni ugao
kamere?



Razni izazovi

Varijacija unutar jedne klase



Različito osvetljenje



Različite pozadine



Varijacije među klasama/kategorijama

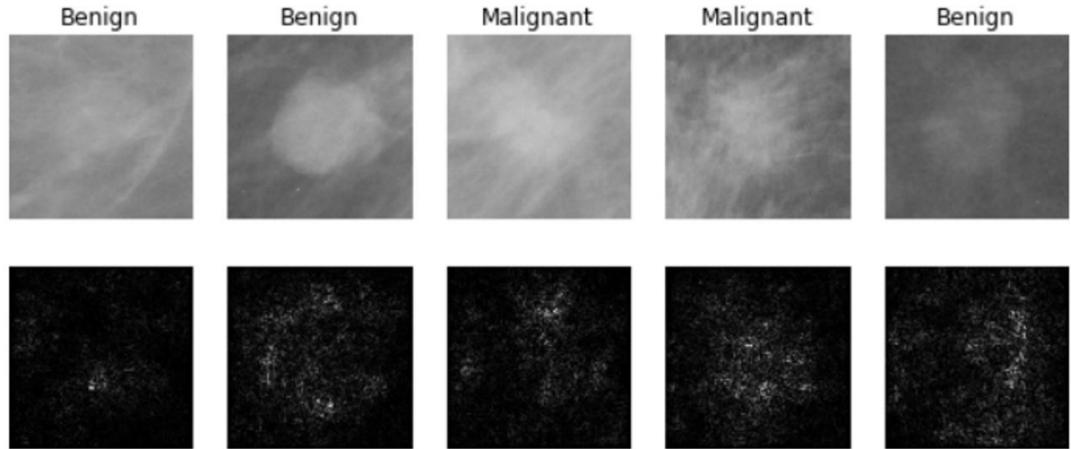
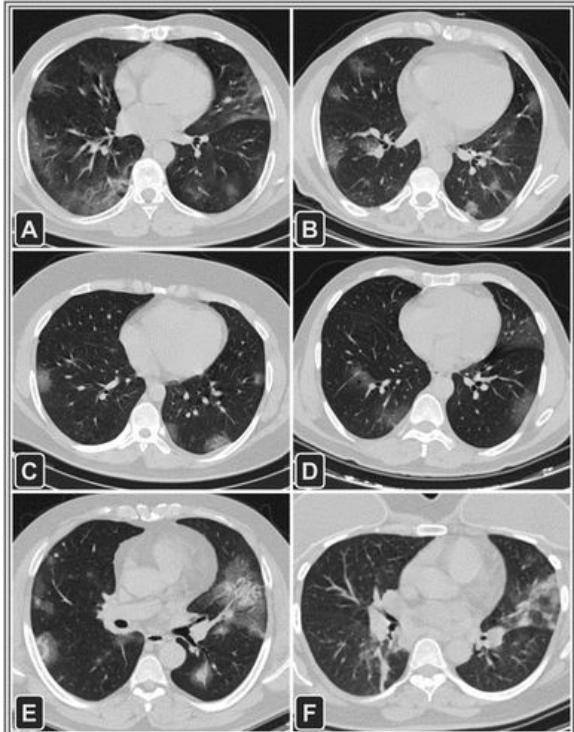


Okluzija

Klasifikator mora biti robustan na različite promjene koje se mogu desiti među slikama.

Upotreba klasifikacije slika

Predviđanja ishoda upale pluća
kod COVID-19 pacijenata



Klasifikacija galaksija



Još jedan primjer: Identifikacija kitova!

The screenshot shows a competition page for the "Humpback Whale Identification Challenge". At the top left is a "Playground Prediction Competition" logo. The main title is "Humpback Whale Identification Challenge" with the subtitle "Can you identify a whale by the picture of its fluke?". Below this is a large image of a whale's tail (fluke) in the water. A blue arrow points from the text "Inside Kaggle you'll find all the code & data you need to do your data science work. Use over 50,000 public datasets and 400,000 public notebooks to conquer any analysis in no time." down to the "Kaggle · 528 teams · 2 years ago" badge. On the right is a "Leaderboard" table with the following data:

#	Δpub	Team Name	Notebook	Team Members	Score	Entries	Last
1	—	Martin Piotte			0.78563	84	2y
2	—	Olga Moskvyak			0.64924	90	2y
3	—	pi-null-mezon			0.63617	103	2y
4	—	bbrandt			0.61281	31	2y
5	—	NTU_r06922086_feijai			0.60866	64	2y

[Overview](#) [Data](#) [Notebooks](#) [Discussion](#) [Leaderboard](#) [Rules](#)

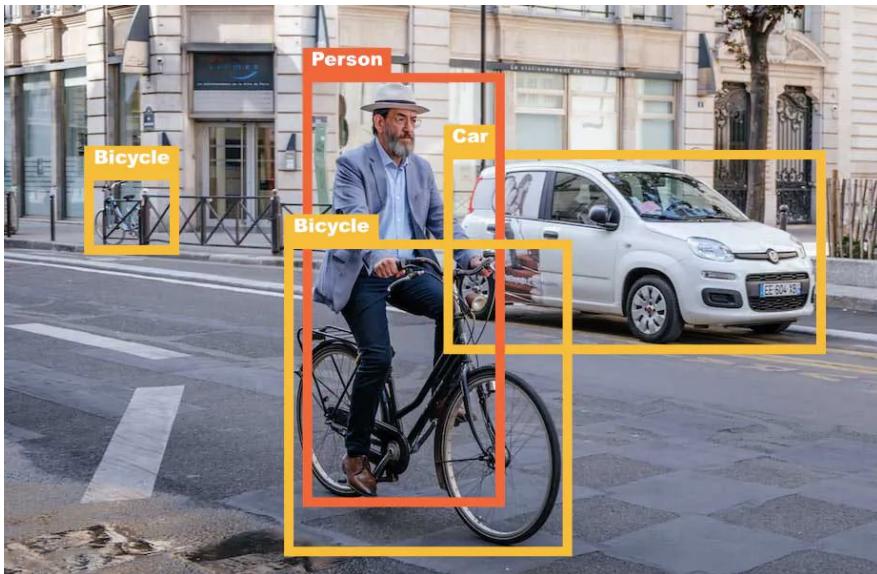
[Join Competition](#)

File descriptions

- train.zip - a folder containing the training images
- train.csv - maps the training Image to the appropriate whale Id . Whales that are not predicted to have a label identified in the training data should be labeled as new_whale .
- test.zip - a folder containing the test images to predict the whale Id
- sample_submission.csv - a sample submission file in the correct format

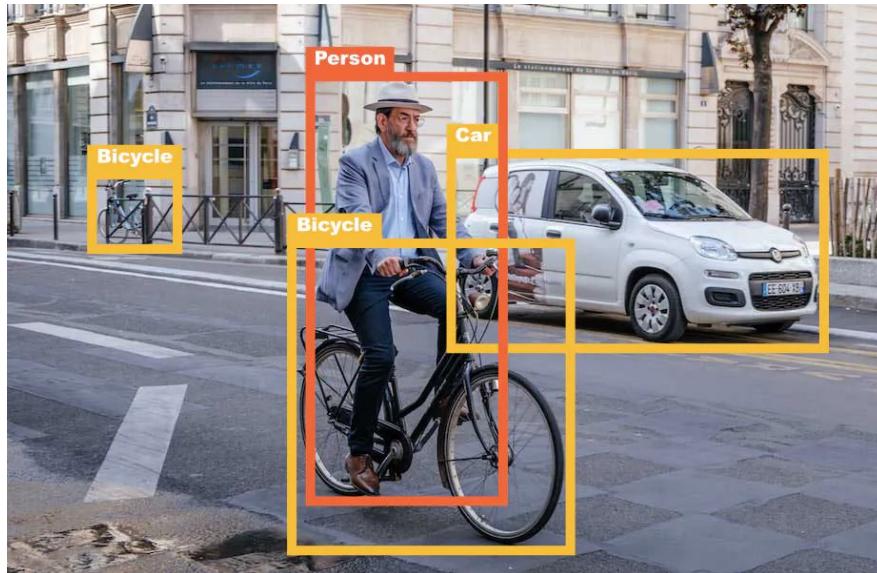
Inside Kaggle you'll find all the code & data you need to do your data science work. Use over 50,000 public datasets and 400,000 public notebooks to conquer any analysis in no time.

Detekcija objekata



→ **automobil**
biciklo
vozi
osoba
zgrada
ulica
...

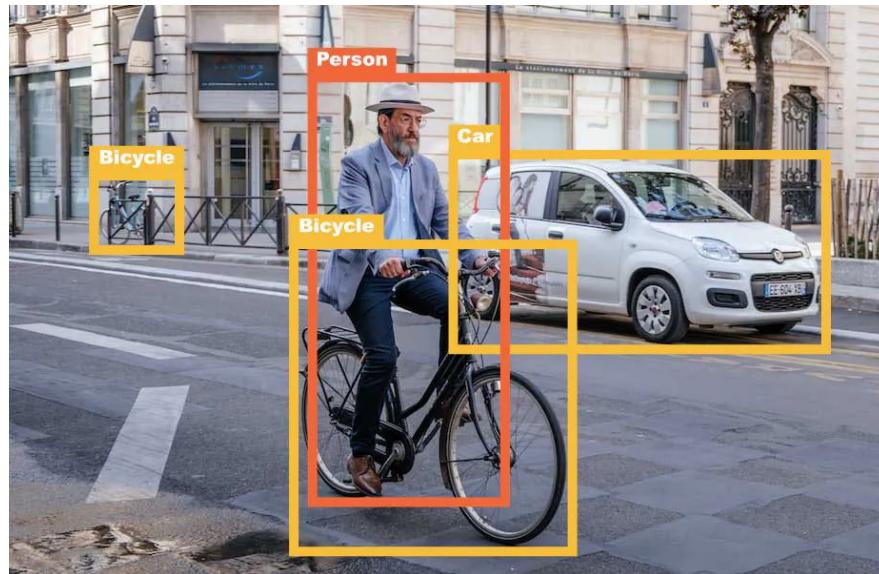
Titlovanje slika / Dodjeljivanje naslova



automobil
biciklo
vozi
osoba
zgrada
ulica
...
<STOP>

Osoba ...

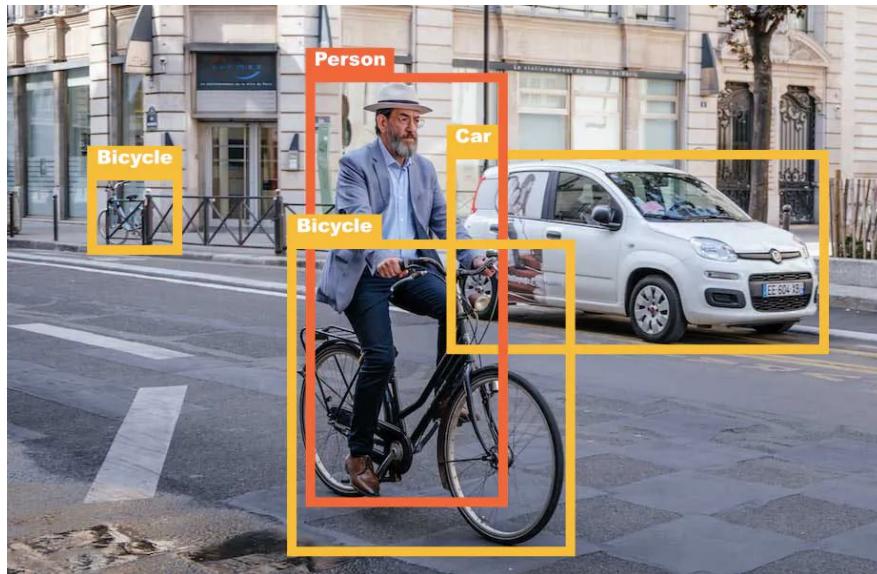
Titlovanje slika / Dodjeljivanje naslova



automobil
biciklo
vozi
osoba
zgrada
ulica
...
<STOP>

Osoba vozi...

Titlovanje slika / Dodjeljivanje naslova



automobil
biciklo
vozi
osoba
zgrada
ulica
...
<STOP>

Osoba vozi biciklo...

Igrice: koji sljedeći potez da napravim?



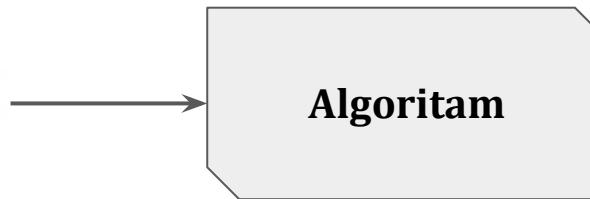
(2, 19)
(4, 19)
...
→ (1, 19)
...
(14, 8)
...

Sljedeći potez?

Klasifikator slika



Ulaz



Izlaz: dodijeli neku predefinisanu kategoriju slici na ulazu

Pas
Mačka
Ptica
Patka
Kamion

```
def classify_image(image):  
    # some magic here  
    return class_label
```

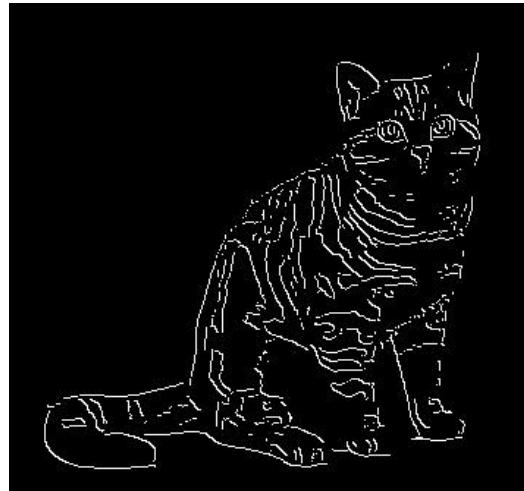
Ne postoji algoritam koji može prepoznati "mačku" na osnovu brojeva. Problem nije jednostavan.

Klasifikator slika

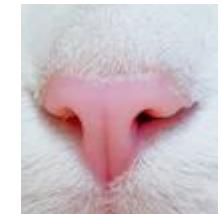
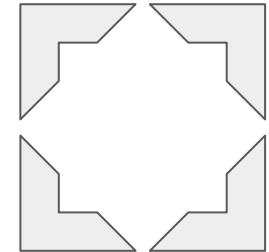
Ulaz



Pronađi ivice



Pronađi neke
značajke



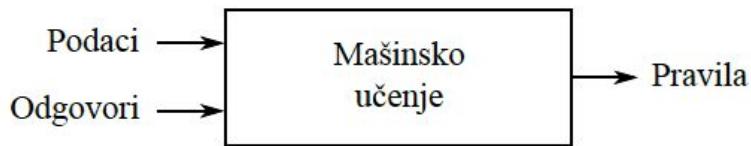
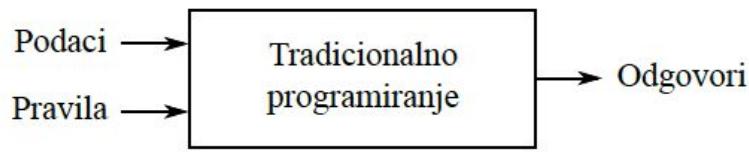
Suština ovakvog pristupa je kodiranje ljudskog znanja
o samom izgledu objekta.

Međutim, problem je što će uvijek biti izuzetaka!



Rješenje: mašinsko učenje

- ❑ Pristup koji je vođen podacima (*data-driven approach*)
- ❑ Ideja: umjesto da ručno kodiramo značajke objekata - hajde da (na)učimo iz podataka.
- ❑ Uobičajeni koraci:
 - ❑ Nađi ili skupi puno podataka i svakom podatku u skupu pridruži labelu.
 - ❑ Koristi jedan od metoda mašinskog učenja da istreniraš klasifikator.
 - ❑ Napravi evaluaciju klasifikatora nad nekim skupom testnih podataka.



Rješenje: mašinsko učenje

- ❑ Nova paradigma programiranja: računar se programira preko podataka koje dovodimo na ulaz programa, a ne preko eksplicitnih pravila.
- ❑ Program možemo mijenjati dovođenjem novih podataka!
- ❑ Ova paradigma je posebno važna za probleme za koje ne znam kako riješiti ili pristupiti njihovom rješavanju.

```
def train(image):  
    # machine learning  
    return model
```

```
def predict(image):  
    # Use model to predict new data  
    return test_label
```

Avion



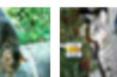
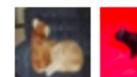
Automobil



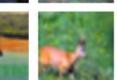
Ptica



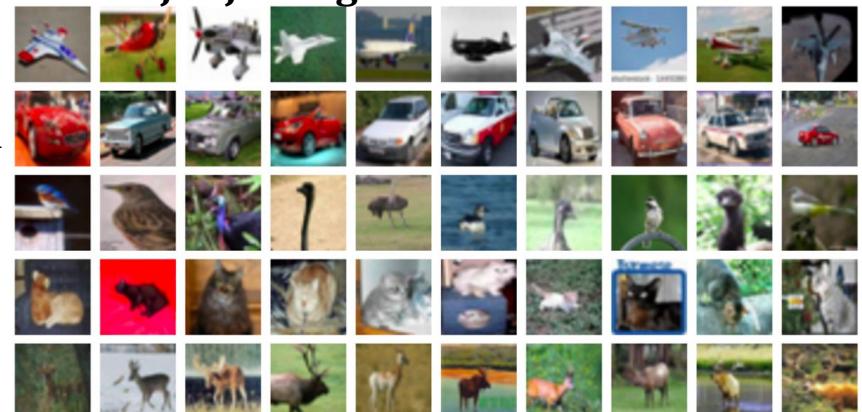
Mačka



Srna



Primjer jednog dataset-a



Model mašinskog učenja će pokušati da nauči statističku povezanost između podataka (ulaznih slika i labela).

Poznati izvori podataka: MNIST



10 klasa: Cifre od 0 do 9
28x28 slike sive nijanse
50k slika za treniranje
10k slika za testiranje

<http://yann.lecun.com/exdb/mnist/>

Poznati izvori podataka: MNIST

CLASSIFIER	PREPROCESSING	TEST ERROR RATE (%)	Reference
Linear Classifiers			
linear classifier (1-layer NN)	none	12.0	LeCun et al. 1998
linear classifier (1-layer NN)	deskewing	8.4	LeCun et al. 1998
pairwise linear classifier	deskewing	7.6	LeCun et al. 1998
K-Nearest Neighbors			
K-nearest-neighbors, Euclidean (L2)	none	5.0	LeCun et al. 1998
K-nearest-neighbors, Euclidean (L2)	none	3.09	Kenneth Wilder, U. Chicago
K-nearest-neighbors, L3	none	2.83	Kenneth Wilder, U. Chicago
K-nearest-neighbors, Euclidean (L2)	deskewing	2.4	LeCun et al. 1998
K-nearest-neighbors, Euclidean (L2)	deskewing, noise removal, blurring	1.80	Kenneth Wilder, U. Chicago
K-nearest-neighbors, L3	deskewing, noise removal, blurring	1.73	Kenneth Wilder, U. Chicago
K-nearest-neighbors, L3	deskewing, noise removal, blurring, 1 pixel shift	1.33	Kenneth Wilder, U. Chicago
K-nearest-neighbors, L3	deskewing, noise removal, blurring, 2 pixel shift	1.22	Kenneth Wilder, U. Chicago

<http://yann.lecun.com/exdb/mnist/>

Non-Linear Classifiers			
40 PCA + quadratic classifier	none	3.3	LeCun et al. 1998
1000 RBF + linear classifier	none	3.6	LeCun et al. 1998
SVMs			
SVM, Gaussian Kernel	none	1.4	
SVM deg 4 polynomial	deskewing	1.1	LeCun et al. 1998
Reduced Set SVM deg 5 polynomial	deskewing	1.0	LeCun et al. 1998
Virtual SVM deg-9 poly [distortions]	none	0.8	LeCun et al. 1998
Virtual SVM, deg-9 poly, 1-pixel jittered	none	0.68	DeCoste and Scholkopf, MLJ 2002
Virtual SVM, deg-9 poly, 1-pixel jittered	deskewing	0.68	DeCoste and Scholkopf, MLJ 2002
Virtual SVM, deg-9 poly, 2-pixel jittered	deskewing	0.56	DeCoste and Scholkopf, MLJ 2002
Convolutional nets			
Convolutional net LeNet-1	subsampling to 16x16 pixels	1.7	LeCun et al. 1998
Convolutional net LeNet-4	none	1.1	LeCun et al. 1998
Convolutional net LeNet-4 with K-NN instead of last layer	none	1.1	LeCun et al. 1998
Convolutional net LeNet-4 with local learning instead of last layer	none	1.1	LeCun et al. 1998
large/deep conv. net, 1-20-40-60-80-100-120-120-10 [elastic distortions]	none	0.35	Ciresan et al. IJCAI 2011
committee of 7 conv. net, 1-20-P-40-P-150-10 [elastic distortions]	width normalization	0.27 +0.02	Ciresan et al. ICDAR 2011
committee of 35 conv. net, 1-20-P-40-P-150-10 [elastic distortions]	width normalization	0.23	Ciresan et al. CVPR 2012

Poznati izvori podataka: CIFAR10

Avion



10 klasa

Automobil



50k slika za treniranje (5k po klasi)

Ptica



10k slika za testiranje (1k po klasi)

Mačka



32x32 RGB slika

Srna



Pas



Žaba



Konj



Brod



Kamion



<https://www.kaggle.com/c/cifar-10/overview>

#	Δpub	Team Name	Team Members	Score	Entries	Last
1	—	DeepCNet		0.95530	18	6y
2	—	jiki		0.94740	42	6y
3	—	Anil Thomas		0.94300	3	6y
4	—	Frank Shar		0.94190	13	6y
5	—	nagadomi		0.94150	16	6y

Poznati izvori podataka: CIFAR100

- Ovaj skup podataka sličan je CIFAR-10, osim što ima 100 klasa koje sadrže po 600 slika. Postoji 500 slika za trening i 100 slika za testiranje po klasi. 100 klasa u CIFAR-100 grupirano je u 20 super klasa. Svaka slika dolazi s oznakom "fine" (klasa kojoj pripada) i oznakom "coarse" (superklasa kojoj pripada).

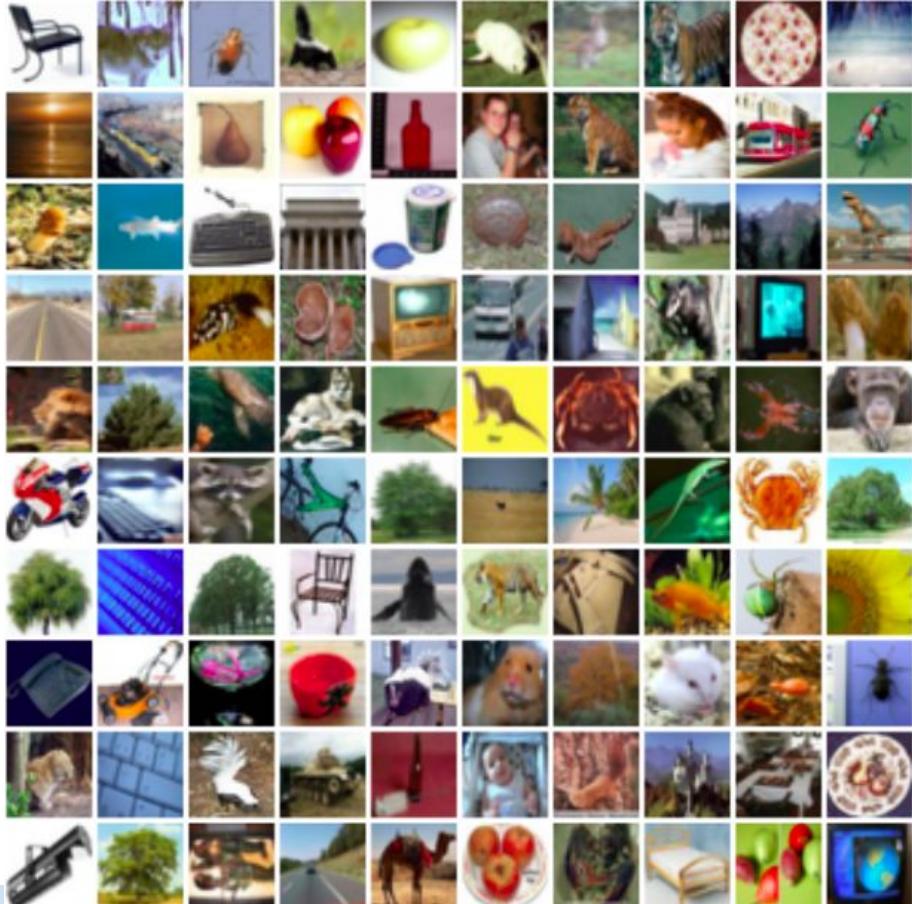
Superclass

aquatic mammals
fish
flowers
food containers
fruit and vegetables
household electrical devices
household furniture
insects
large carnivores
large man-made outdoor things
large natural outdoor scenes
large omnivores and herbivores
medium-sized mammals
non-insect invertebrates
people
reptiles
small mammals
trees
vehicles 1
vehicles 2

Classes

beaver, dolphin, otter, seal, whale
aquarium fish, flatfish, ray, shark, trout
orchids, poppies, roses, sunflowers, tulips
bottles, bowls, cans, cups, plates
apples, mushrooms, oranges, pears, sweet peppers
clock, computer keyboard, lamp, telephone, television
bed, chair, couch, table, wardrobe
bee, beetle, butterfly, caterpillar, cockroach
bear, leopard, lion, tiger, wolf
bridge, castle, house, road, skyscraper
cloud, forest, mountain, plain, sea
camel, cattle, chimpanzee, elephant, kangaroo
fox, porcupine, possum, raccoon, skunk
crab, lobster, snail, spider, worm
baby, boy, girl, man, woman
crocodile, dinosaur, lizard, snake, turtle
hamster, mouse, rabbit, shrew, squirrel
maple, oak, palm, pine, willow
bicycle, bus, motorcycle, pickup truck, train
lawn-mower, rocket, streetcar, tank, tractor

Poznati izvori podataka: CIFAR100



100 klasa

50k slika za treniranje (500 po klasi)

10k slika za testiranje (100 po klasi)

32x32 RGB slika

20 super klasa u kojoj je po 5 klasa

<https://benchmarks.ai/cifar-100>

CIFAR-100

Classify 32x32 colour images into 100 categories.

Method	(expand all collapse all)	Accuracy (%)
Big Transfer (BiT): General Visual Representation Learning	(Dec 2019)	93.51%
Alexander Kolesnikov, Lucas Beyer, Xiaohua Zhai, Joan Puigcerver, Jessica Yung, Sylvain Gelly, Neil Houlsby		
Transfer of pre-trained representations improves sample efficiency and simplifies hyperparameter tuning when training deep neural networks for vision. We revisit the paradigm of pre-training on large supervised datasets and fine-tuning the model on a target task. We scale up pre-training, and propose a simple recipe that we call Big Transfer (BiT). By combining a few carefully selected components, and transferring using a simple heuristic, we achieve strong performance on over 20 datasets. BiT performs well across a surprisingly wide range of data regimes -- from 1 example per class to 1M total examples. BiT achieves 87.5% top-1 accuracy on ILSVRC-2012, 99.4% on CIFAR-10, and 76.3% on the 19 task Visual Task Adaptation Benchmark (VTAB). On small datasets, BiT attains 76.8% on ILSVRC-2012 with 10 examples per class, and 97.0% on CIFAR-10 with 10 examples per class. We conduct detailed analysis of the main components that lead to high transfer performance.		
EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks	(May 2019, arXiv 2019)	91.70%
Mingxing Tan, Quoc V. Le		
Convolutional Neural Networks (ConvNets) are commonly developed at a fixed resource budget, and then scaled up for better accuracy if more resources are available. In this paper, we systematically study model scaling and identify that carefully balancing network depth, width, and resolution can lead to better performance. Based on this observation, we propose a new scaling method that uniformly scales all dimensions of depth/width/resolution using a simple yet highly effective compound coefficient. We demonstrate the effectiveness of this method on scaling up MobileNets and ResNet. To go even further, we use neural architecture search to design a new baseline network and scale it up to obtain a family of models, called EfficientNets, which achieve much better accuracy and efficiency than previous ConvNets. In particular, our EfficientNet-B7 achieves state-of-the-art 84.4% top-1 / 97.1% top-5 accuracy on ImageNet, while being 8.4x smaller and 6.1x faster on inference than the best existing ConvNet. Our EfficientNets also transfer well and achieve state-of-the-art accuracy on CIFAR-100 (91.7%), Flowers (98.8%), and 3 other transfer learning datasets, with an order of magnitude fewer parameters. Source code is at https://github.com/tensorflow/tpu/tree/master/models/official/efficientnet .		

Poznati izvori podataka: ImageNet

Kupola



1000 klasa

Ukupno slika u skupu podataka: 14,197,122
(okt 2020)

~1.3M slika za treniranje (~1.3K po klasi)

50K slika za validaciju (50 po klasi)

100K slika za testiranje (100 po klasi)

Slike su različite veličine, ali se najčešće smanjuju na 256x256 u procesu treniranja.

Performanse se mjeru sa: **Top 5 accuracy**
(Algoritam vrši predikciju o 5 klasa za svaku sliku; jedna od tih pet slika mora biti tačna.)

<http://image-net.org/index>

MIT Places

365 klasa različitih vrsta scena



~8M slika za treniranje

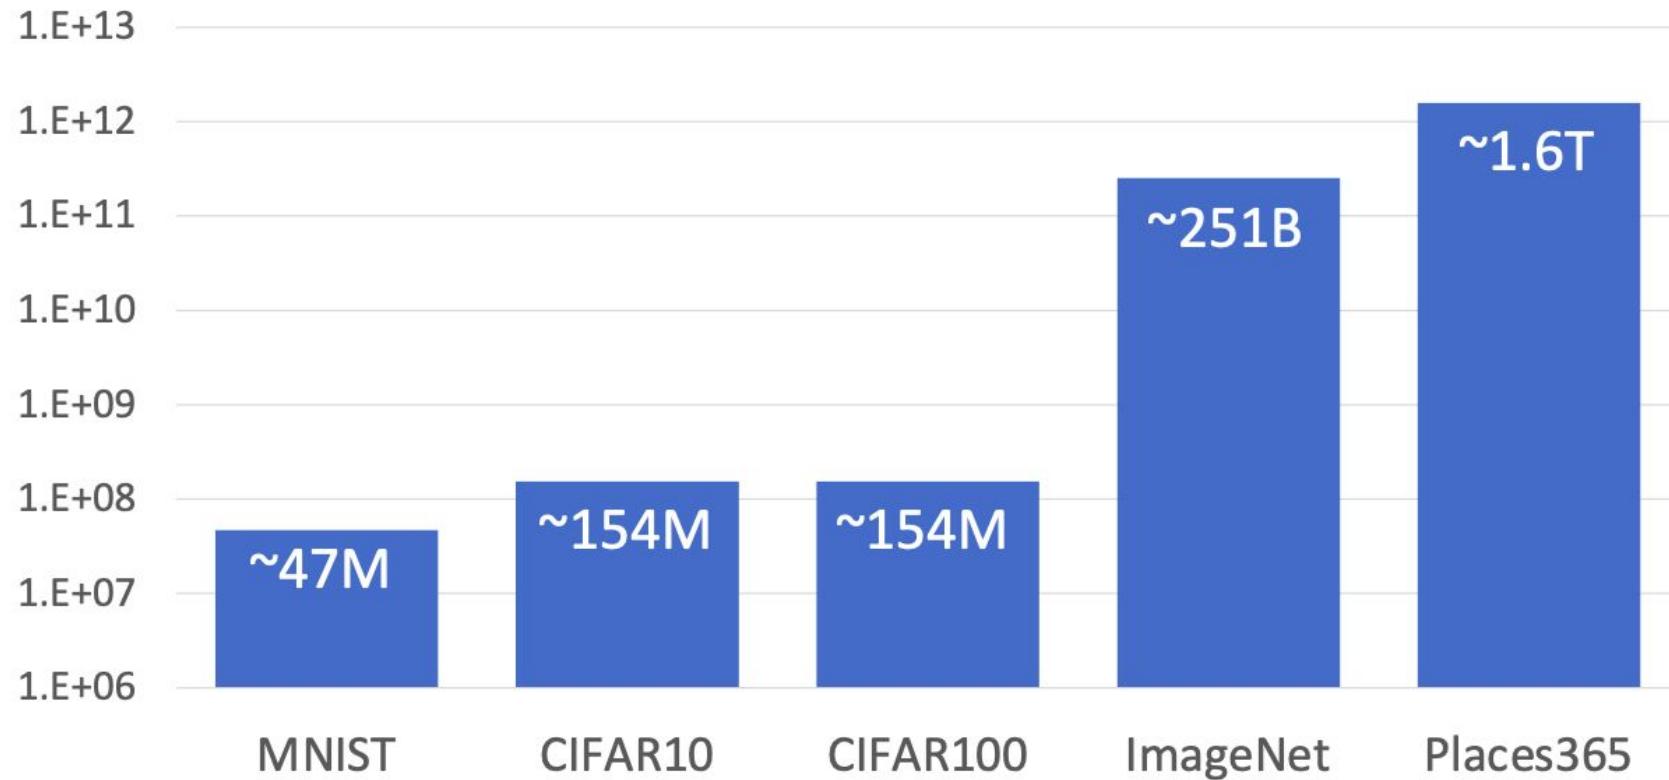
18.25K slika za validaciju (50 po klasi)

328.5K slika za testiranje (900 po klasi)

Slike su različite veličine, ali se najčešće smanjuju na 256x256 u procesu treniranja.

<http://places.csail.mit.edu/>

Poređenje skupova podataka po broju piksela



Prvi klasifikator: Nearest Neighbor

```
def train(image):  
    # machine learning  
    return model
```

Memoriziraj sve podatke i labele!

```
def predict(image):  
    # Use model to predict new data  
    return test_label
```

Napravi predikciju na osnovu najslučnije slike korištene u procesu treniranja!

Kako izračunati sličnost između dvije slike?

Metrika za distancu

L1 ili Manhattan distanca: $d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$

test image			
56	32	10	18
90	23	128	133
24	26	178	200
2	0	255	220

-

training image			
10	20	24	17
8	10	89	100
12	16	178	170
4	32	233	112

=

pixel-wise absolute value differences			
46	12	14	1
82	13	39	33
12	10	0	30
2	32	22	108

add → 456

Prvi klasifikator: Nearest Neighbor

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Memoriziraj podatke za treniranje

Prvi klasifikator: Nearest Neighbor

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Za svaki testni podatak:

- Pronađi najsličniju sliku iz skupa za treniranje
- Vrati labelu te slike

Prvi klasifikator: Nearest Neighbor

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Ako imamo N slika u skupu za treniranje, koliko je brzo treniranje?

O(1)

Koliko je brzo testiranje?

O(N)

U praksi se sporo treniranje toleriše, ali ne i sporo testiranje.

Postoji mnogo metoda koje efikasnije mogu da izvrše klasifikaciju:

<https://github.com/facebookresearch/faiss>

Prvi klasifikator: Nearest Neighbor



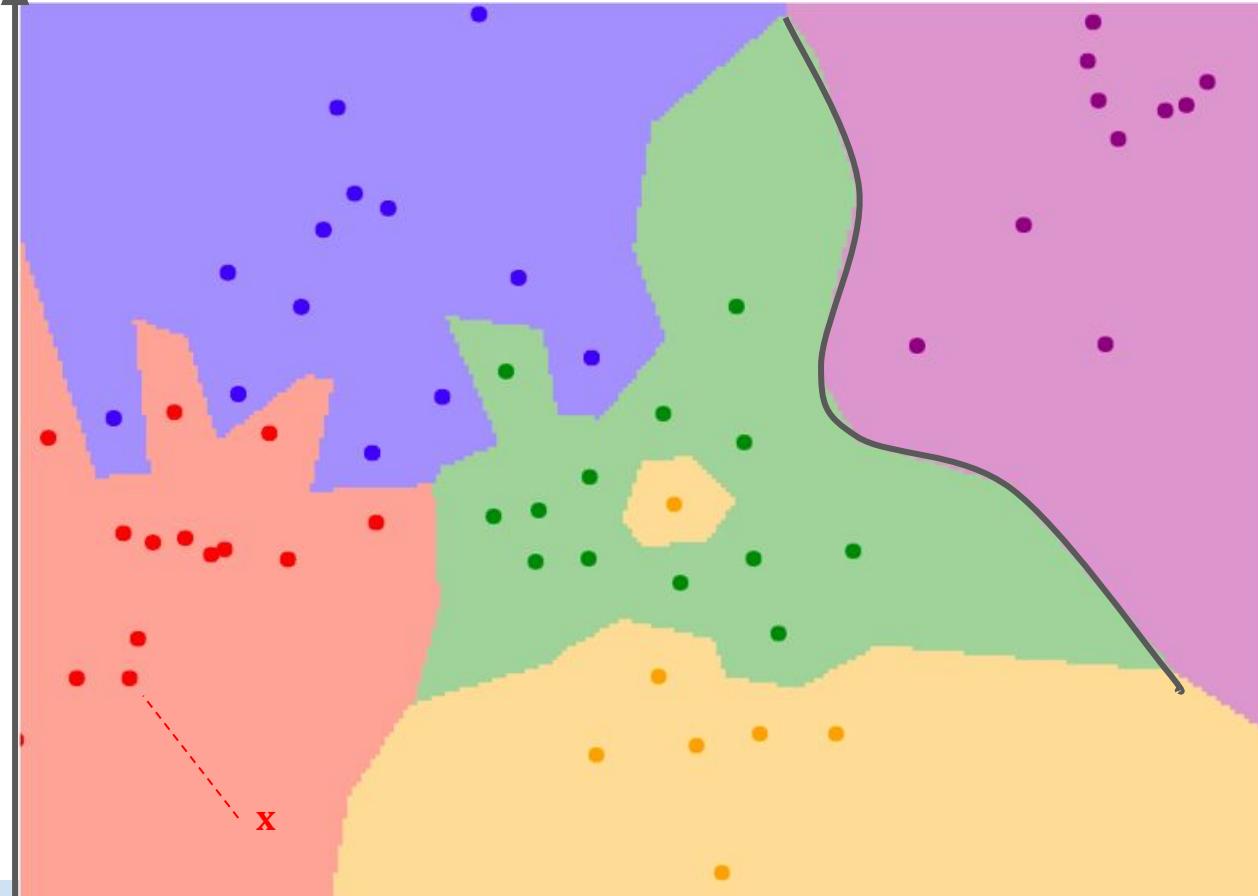
Nedostaci: Nearest Neighbor

Tačke predstavljaju podatke iz skupa za treniranje.

Boja podloge predstavlja klasu.

Outliers

x_1



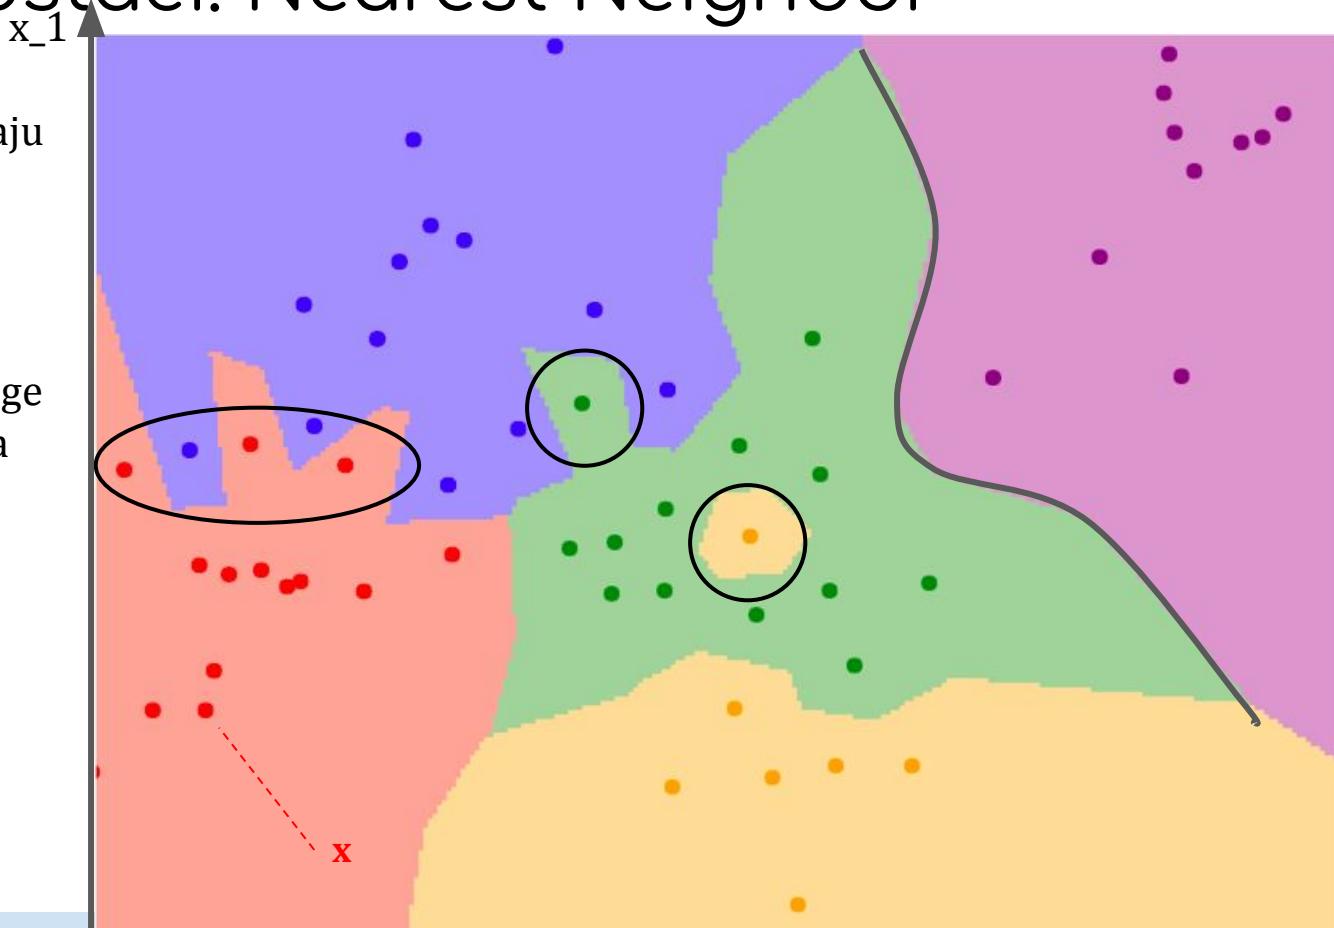
Postoje granice između dva skupa, kao što je prikazano crnom linijom.

Da li je moguće definisati novi algoritam koji će dati preciznije granice?

Nedostaci: Nearest Neighbor

Tačke predstavljaju podatke iz skupa za treniranje.

Boja podloge predstavlja klasu.



Postoje granice između dva skupa, kao što je prikazano crnom linijom.

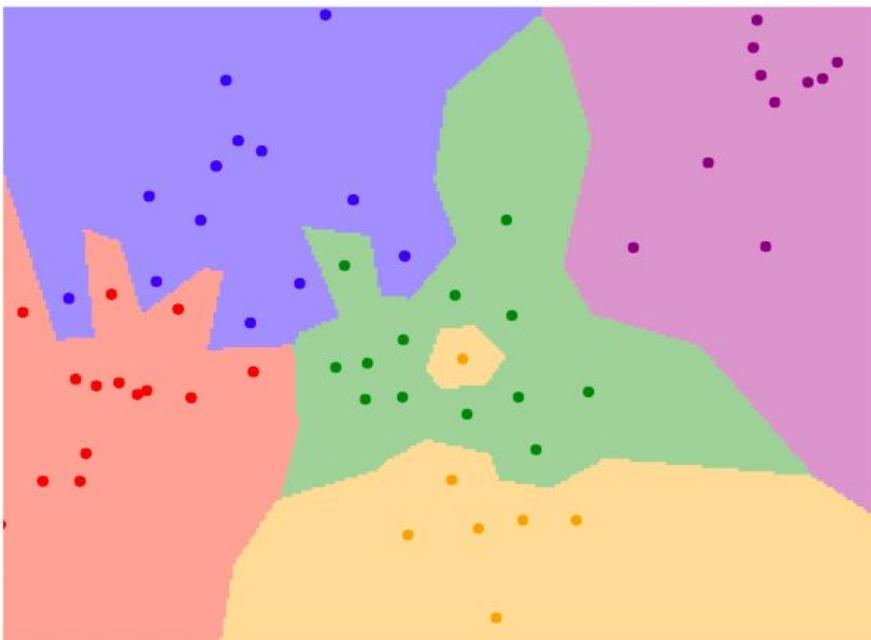
Granice nisu dovoljno precizno određene!

(Outliers)
Da li je moguće definisati novi algoritam koji će dati preciznije granice?

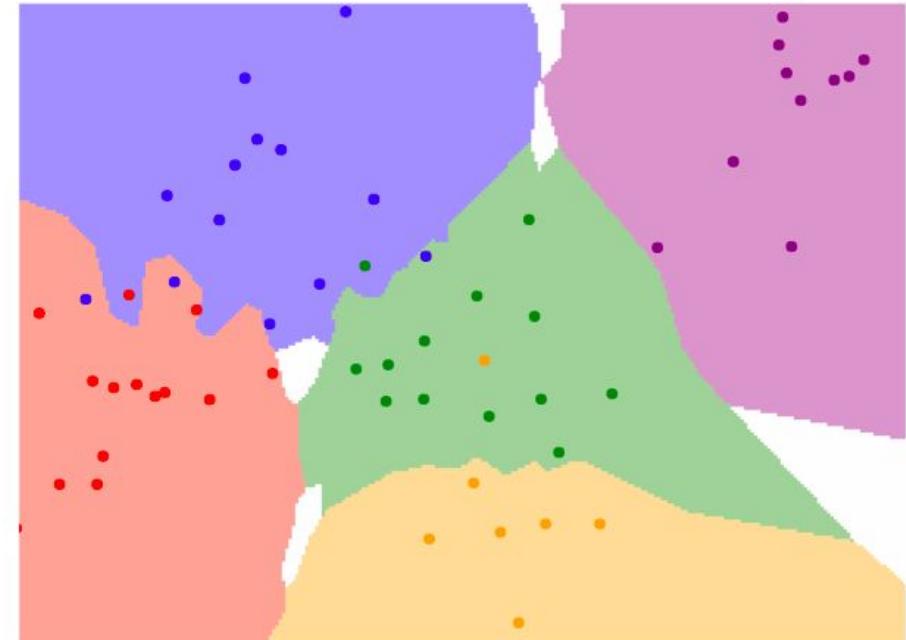
k-Nearest Neighbor

- ❑ Osnovna ideja: Uzeti k najsličnijih slika i “glasati” koju od njih uzeti kao klasu slike koju testiramo.

K = 1



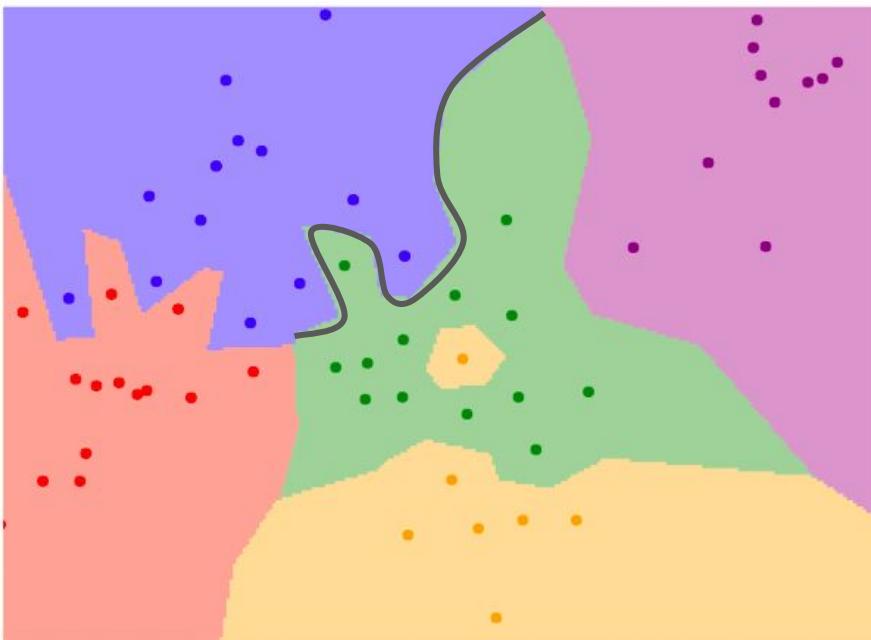
K = 3



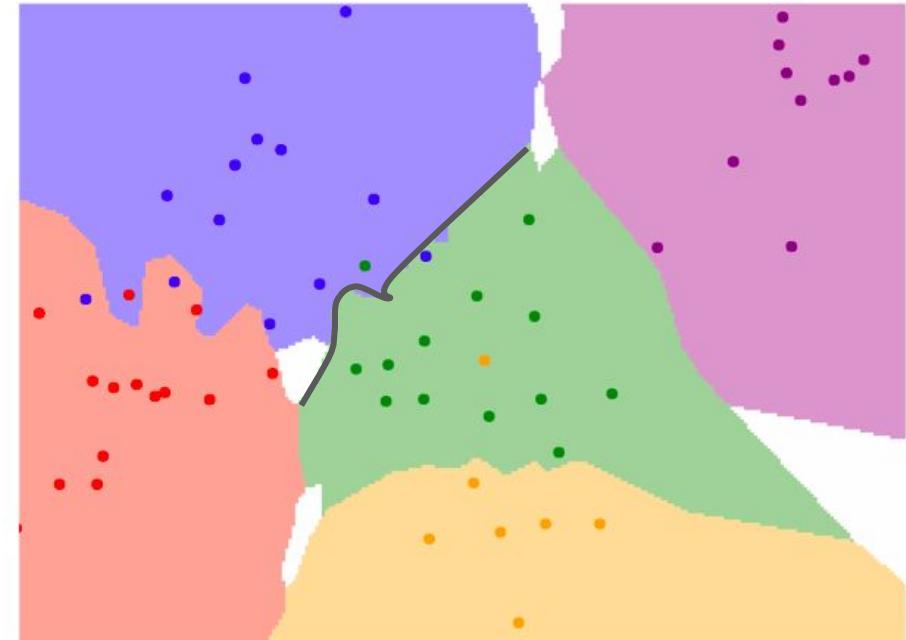
k-Nearest Neighbor

- ❑ Osnovna ideja: Uzeti k najsličnijih slika i “glasati” koju od njih uzeti kao klasu slike koju testiramo.

K = 1



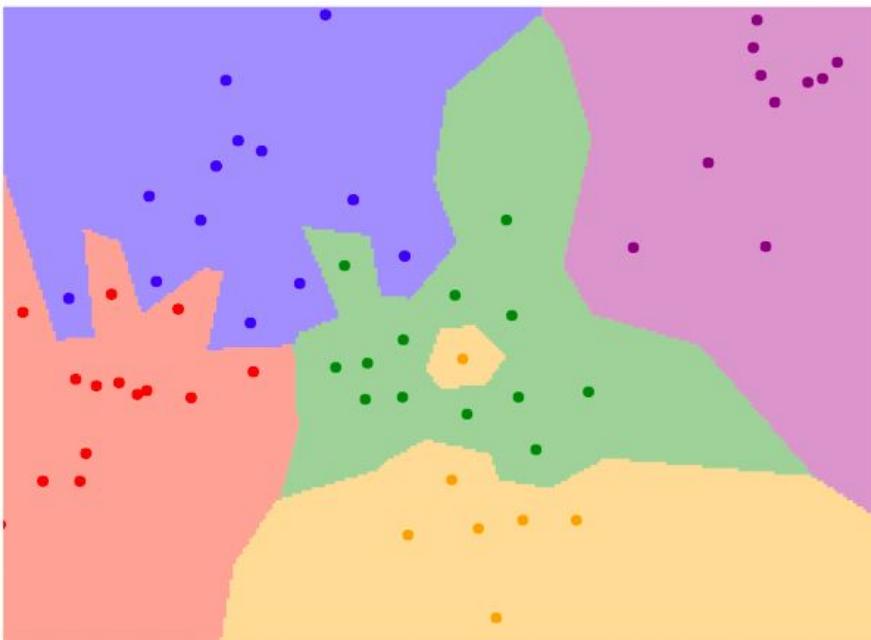
K = 3



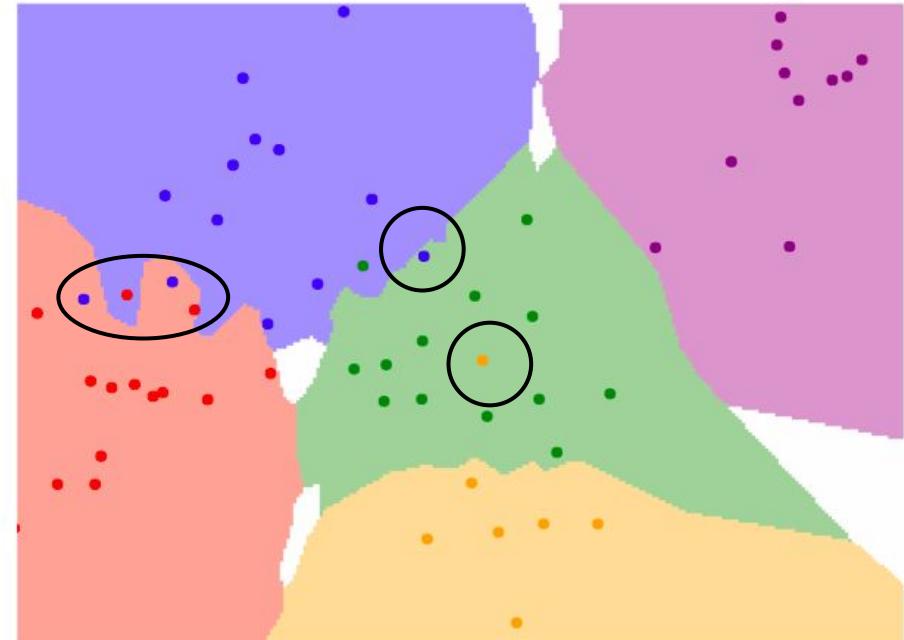
k-Nearest Neighbor

- ❑ Osnovna ideja: Uzeti k najsličnijih slika i “glasati” koju od njih uzeti kao klasu slike koju testiramo.

K = 1



K = 3



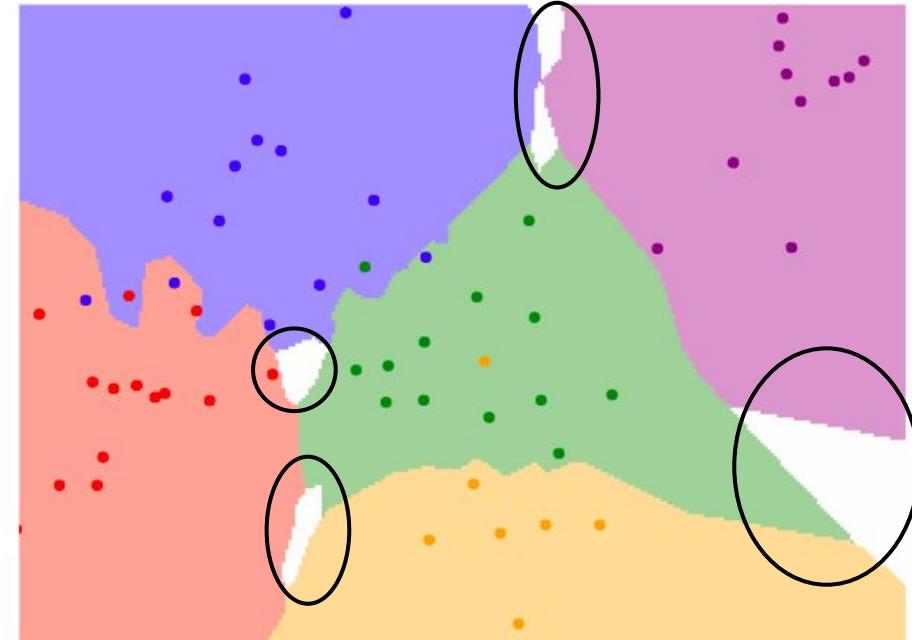
k-Nearest Neighbor

- ❑ Osnovna ideja: Uzeti k najsličnijih slika i “glasati” koju od njih uzeti kao klasu slike koju testiramo.

K = 1



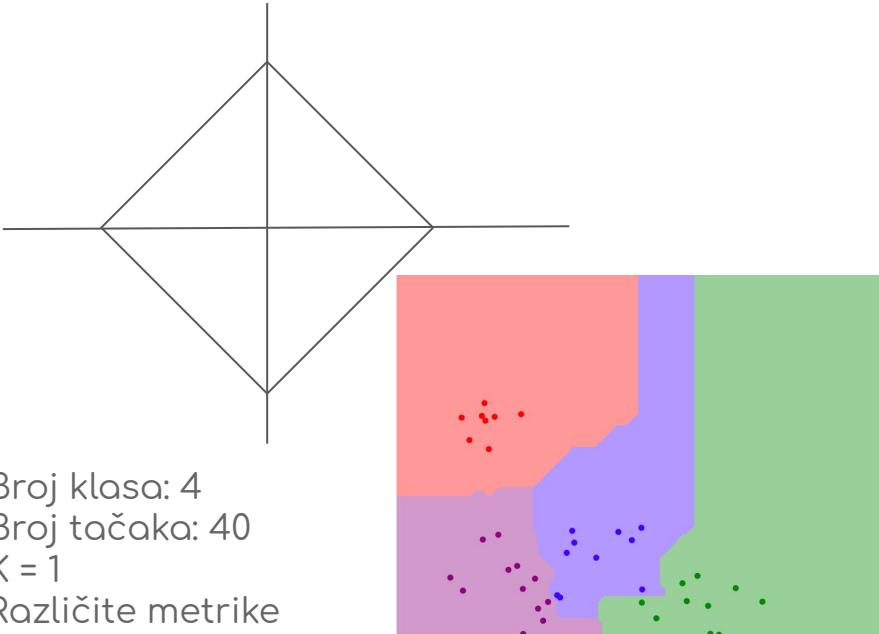
K = 3



Metrika za distancu

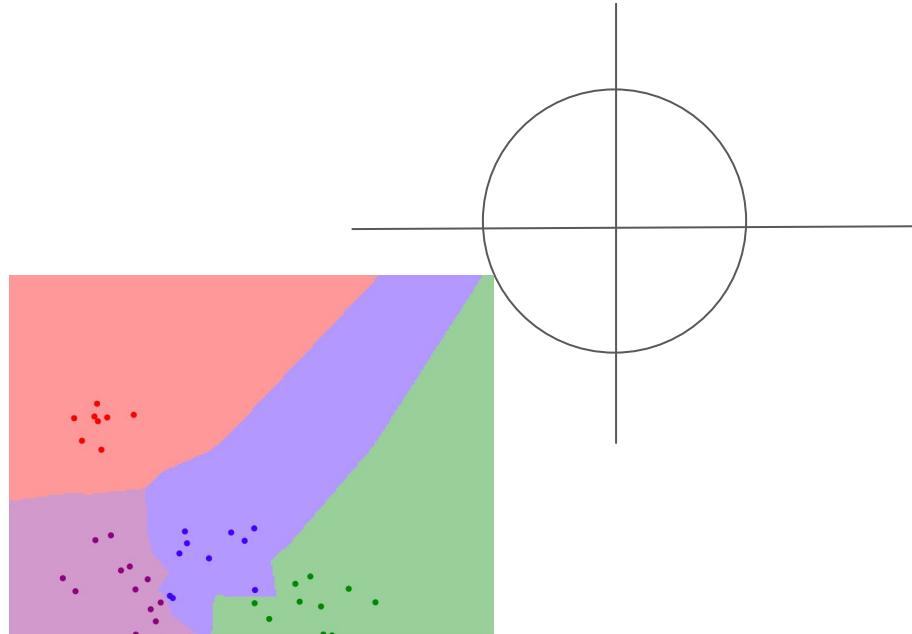
L1 ili *Manhattan* distanca:

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 ili *Euclidean* distanca:

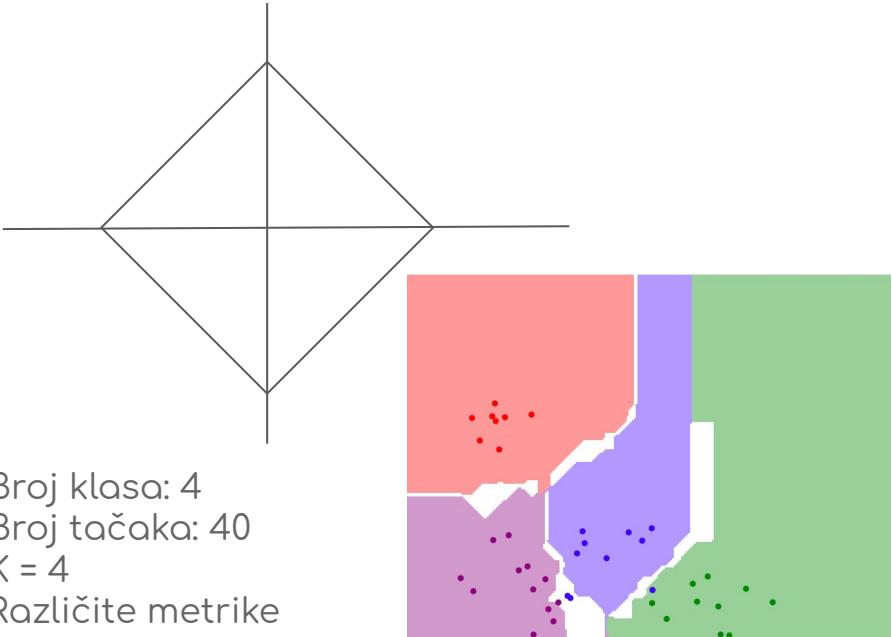
$$d_2(I_1, I_2) = (\sum_p (I_1^p - I_2^p)^2)^{1/2}$$



Metrika za distancu

L1 ili *Manhattan* distanca:

$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



L2 ili *Euclidean* distanca:

$$d_2(I_1, I_2) = (\sum_p (I_1^p - I_2^p)^2)^{1/2}$$

